



Decoupling an Application with Symfony Messenger

ConFoo, Montreal, Canada – February 23th, 2024

© David Buchmann

The background of the image is a wide-angle landscape of the Swiss Alps. In the far distance, majestic mountains with white snow-capped peaks rise against a clear blue sky. Closer to the viewer, there's a dense forest of green and brown trees, suggesting autumn. A grassy hillside with a paved path leads towards the forest. In the bottom left corner, a small wooden structure with a red roof is visible. The overall scene is a mix of natural beauty and human-made infrastructure.

David Buchmann - david@liip.ch
PHP Engineer, Liip SA, Switzerland



We built a data kraken

Product API

massive central index of product data

- Constantly evolving since 2013
- Collecting data from ~40 source systems
Legacy and future sources, always in the process of
migrating to new services
- Cache all data locally in MySQL for quick access
- Aggregate data in Elasticsearch

First approach

- Cronjob looping over all changed products
- Incremental imports: last-modified
- Record command runs and use the timestamp
- Does not scale
- Error during processing: Start over
- Update products many times for different things

Encapsulate Logic

Extract the logic out of commands and controllers.

=> Refactored into Importer classes that can handle a single item at a time.

Decouple with message queue

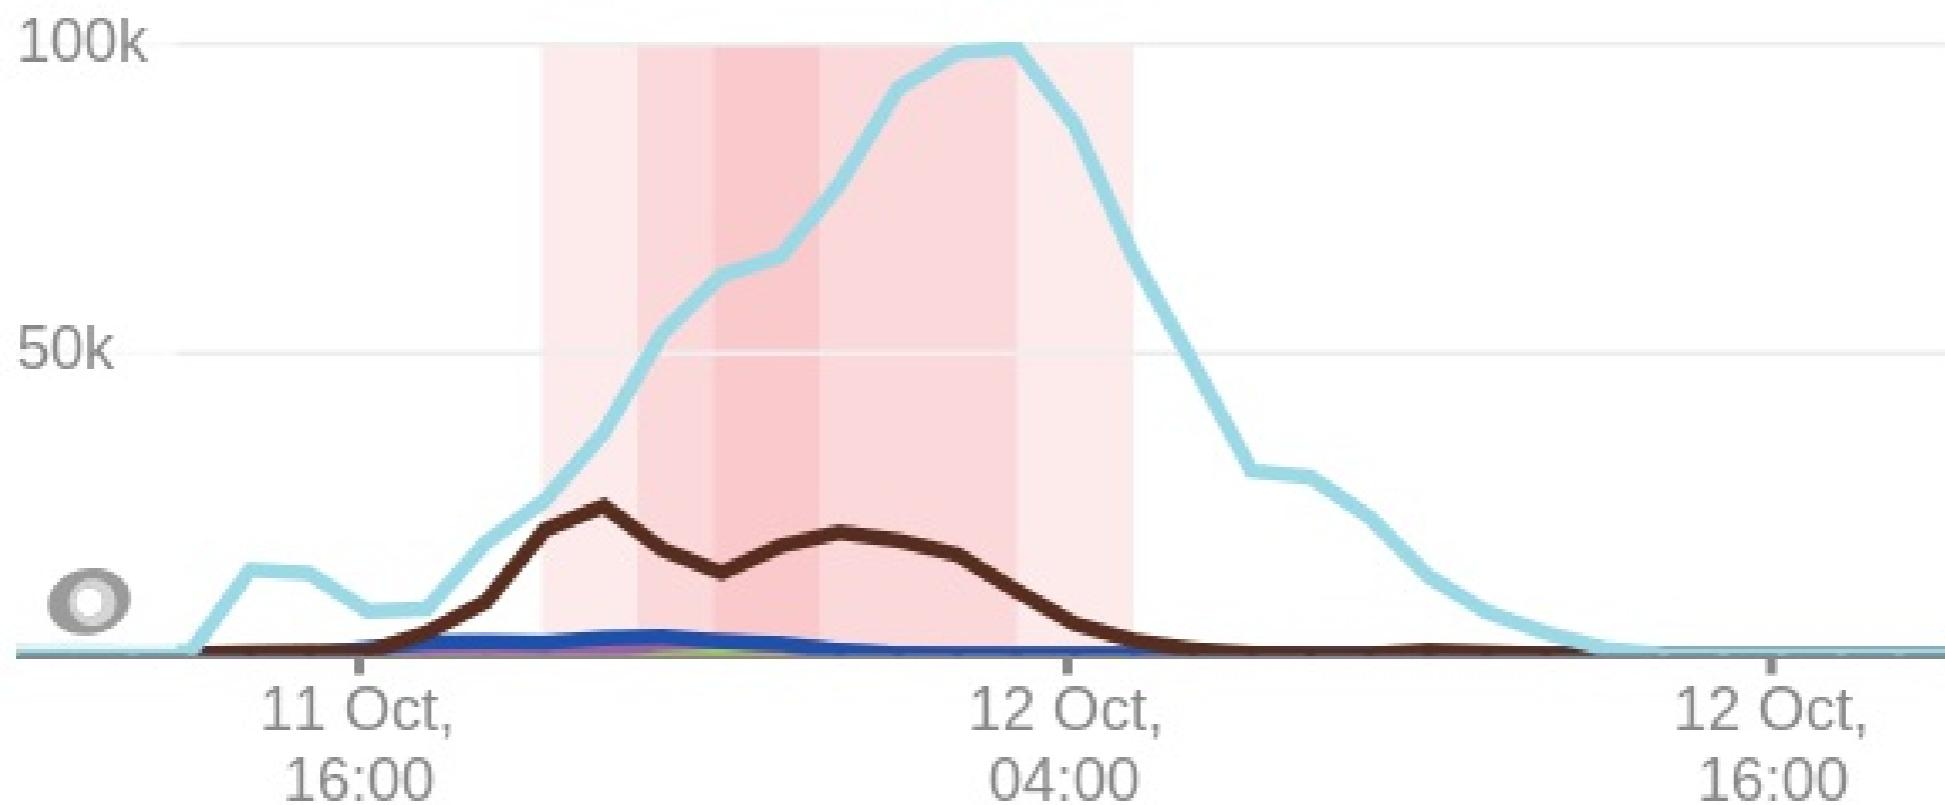
- Importer can be called from command line but also from message processor
- Importing triggers events on data changes
- Event listeners dispatch further messages as needed for next steps
- Perfect if source system can send messages
- Otherwise cronjob to only pull data and store it, then dispatch messages

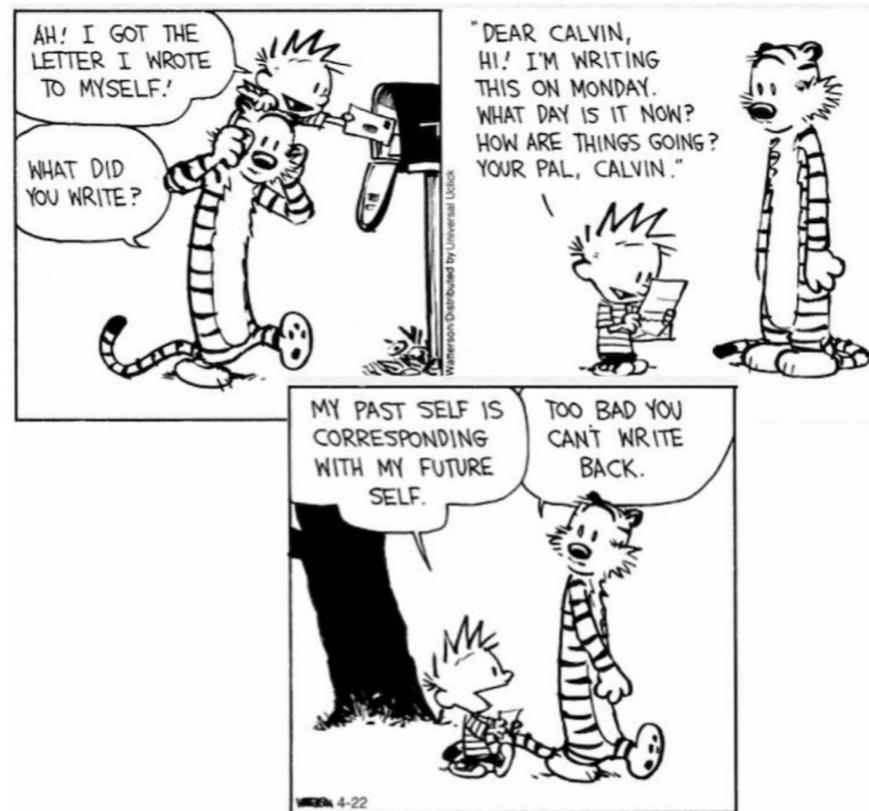
Reaping the benefits

- Parallel processing of data
- Dynamically scale the number of workers
- Failure only affects single items, does not block
 - easy retry
 - abort after several failed attempts

	Messages			Message rates		
State	▼ Ready	Unacked	Total	incoming	deliver / get	ack
running	504,680	99	504,779	2.2/s	31/s	31/s
running	103,692	25	103,717	2.4/s	2.4/s	2.4/s

Message Ready





Messages within the application

- Also split up heavy tasks in the application
- Multiple message-driven steps
- Easy to trigger the next step
- Using Symfony event system to decouple
- Event listeners can dispatch messages: e.g. we have 2 sources of product data, decide based on category which is relevant

e.g. Compile data into Elasticsearch

- SAP data must be processed as fast as possible,
data delivery process is blocking in SAP
- Dispatch message with data as payload
- First consumer stores data locally
- And dispatches an event, which makes another
listener queue this product for reindexing
- Different priorities to process urgent changes fast
(e.g. prices)

Flexible on the outside, consistent on the inside

- Some source systems send messages on changes
- Other systems are polled for changes
- Daily import of CSV files with full data
- Temporal data with time limited validity
- We process everything as message

A large, three-dimensional question mark sign is mounted on a pole against a clear blue sky with a few wispy clouds. The sign has a red outer ring and a white inner ring. A black circle is positioned at the center of the question mark. The text "Any questions so far?" is overlaid in white, bold, sans-serif font across the center of the sign.

**Any questions
so far?**

Demo Time!

github.com/dbu/messenger-demo

<https://md.davidbu.ch>

Neuromancer, William Gibson

Symfony Messenger

Event

Blocking
Can alter behaviour
Deterministic order
Not scalable

Message

Asynchronous
No influence
Undefined order
(of consumers)
Scalable

ConFoo.ca

L//P

Message System Components

- Message (payload application specific)
- Sender (Publisher, Producer)
- Receiver (Subscriber, Consumer)
- Channels (Consume only relevant messages)

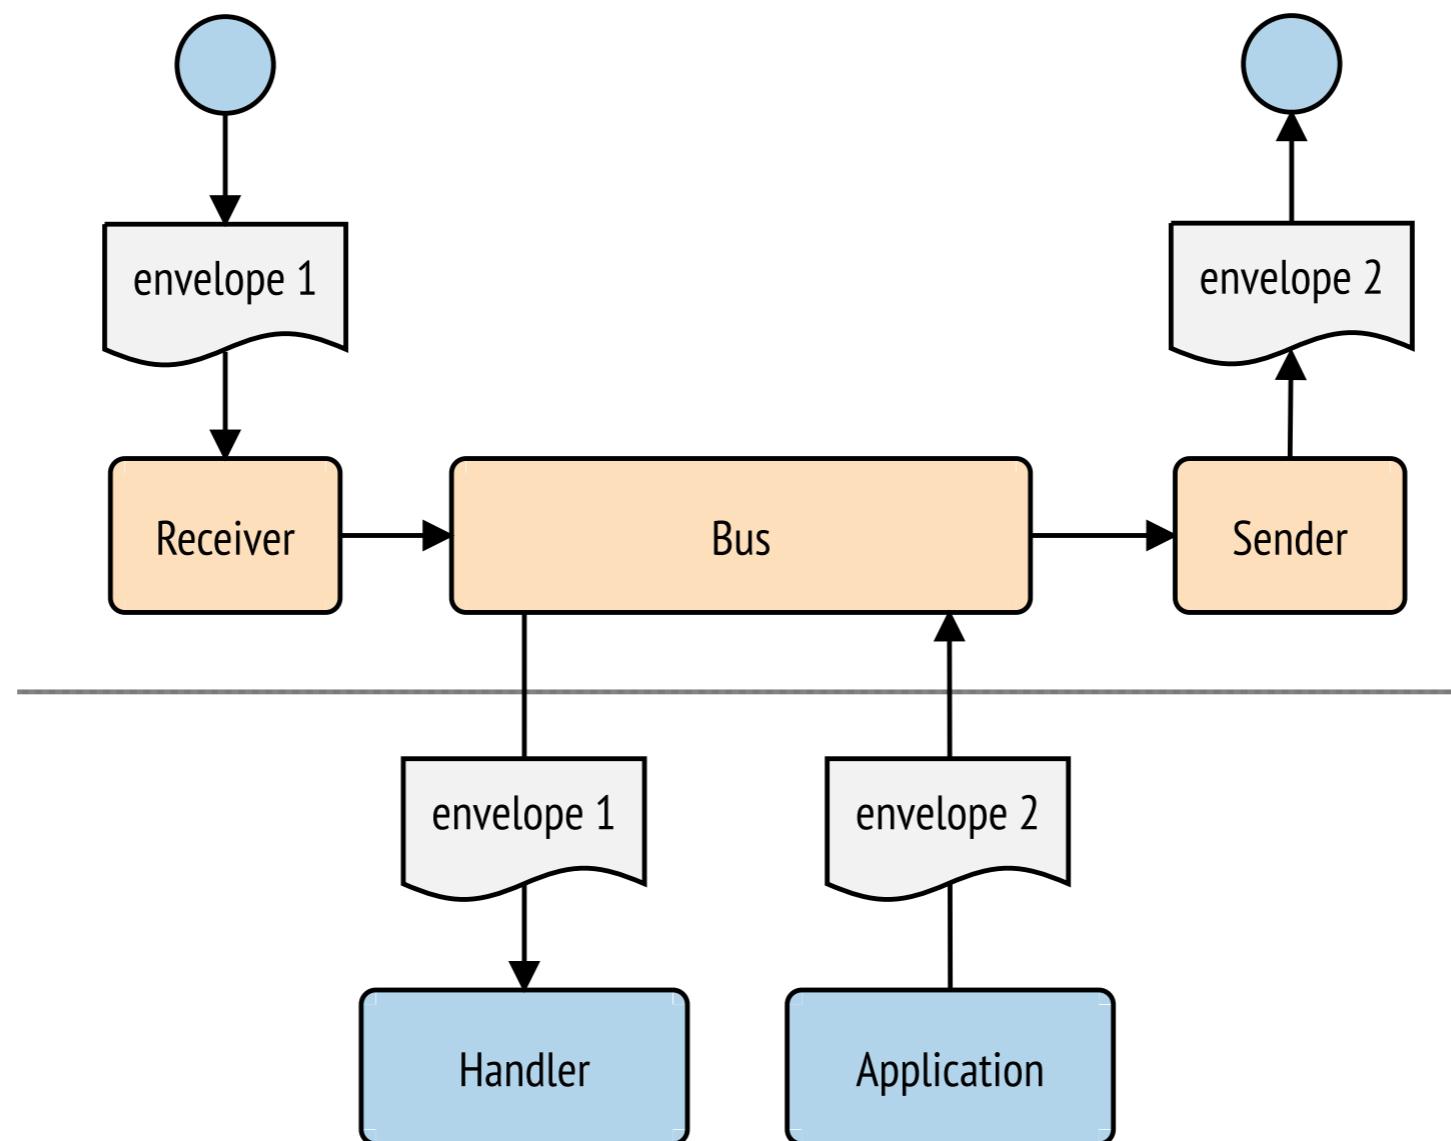
Message System Advantages

- Decoupled (mix systems or languages)
- Scalable (add more producers or consumers)
- Buffering/load management (Async)
- Resilient (buffered when receivers are down)
- Reliable (acknowledgement and retry)
- Order (sequential if only 1 instance)
- Extensions (e.g. message priorities)

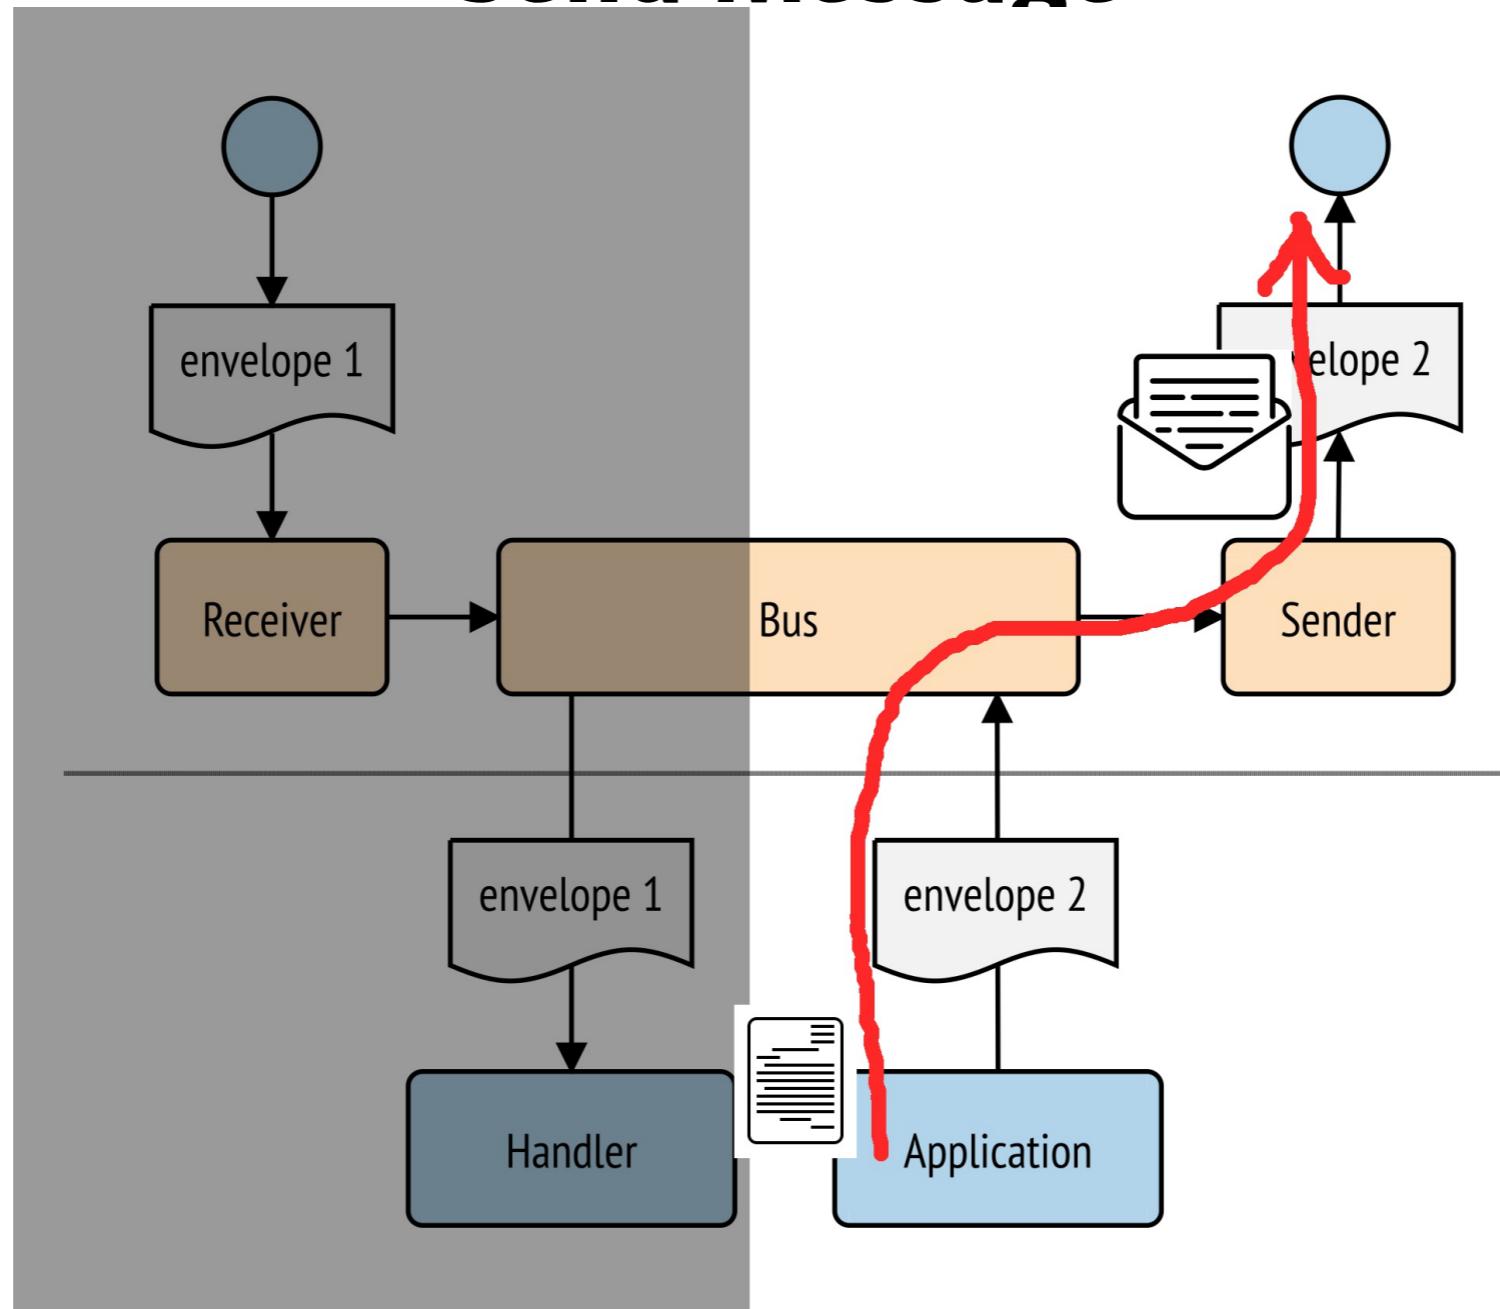
Why Symfony Messenger?

- High level functionality
- Abstract from specific message transport
- Nice and modern architecture

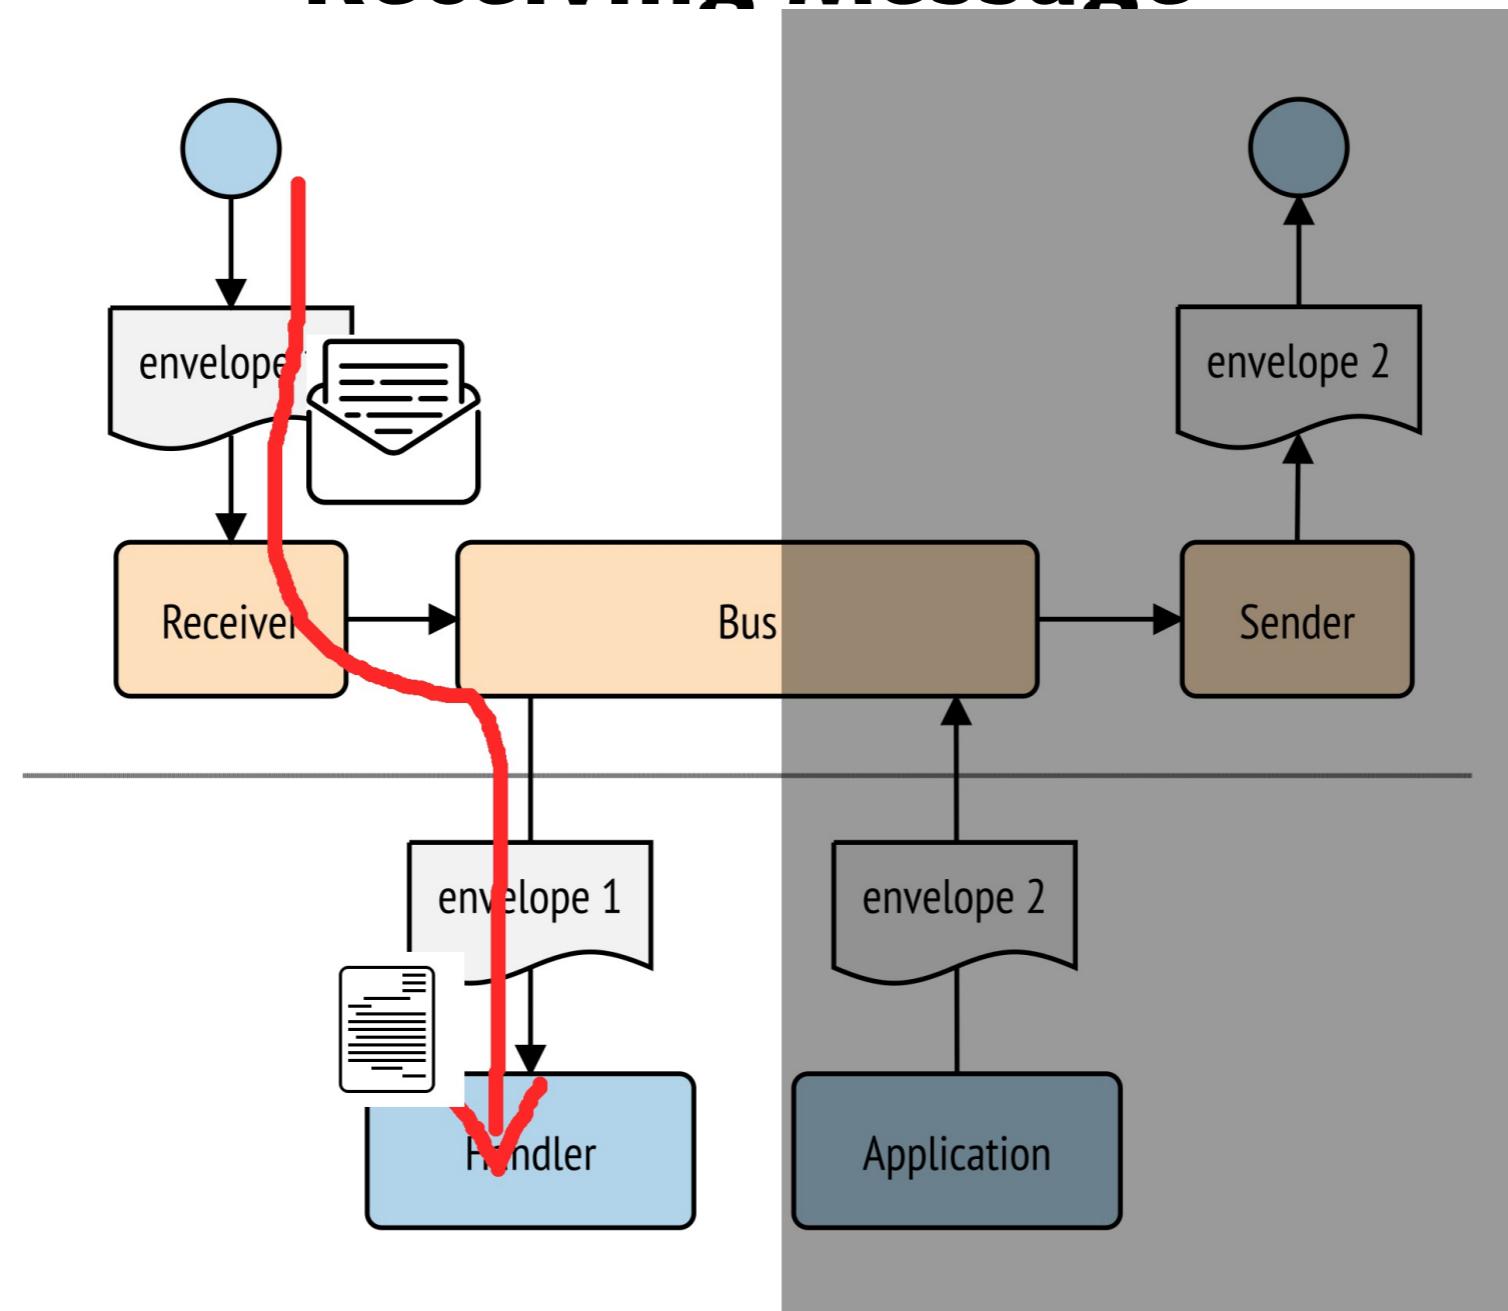
The flow of Symfony Messenger



Send Message



Receiving Message



Message

- This is your class. It is serialized by the messenger.
- Small (Entity ID rather than the full entity data)
- Specific

```
1 class UpdatePromotion
2 {
3     public function __construct(
4         public string $id
5     ) {}
6 }
```

Sending a message

```
1 private MessageBusInterface $bus;  
2 ...  
3 $this->bus->dispatch(new UpdatePromotion($id));
```

Stamps (Metadata)

```
1 $stamps[] = new DelayStamp(20);  
2 $this->bus->dispatch(  
3     new UpdatePromotion($id),  
4     $stamps  
5 );
```

Routing

- What queue for which message?
- Control which workers to scale
- Better monitoring

framework:

 messenger:

 transports:

 neuromancer:

 dsn: '%env(MESSENGER_TRANSPORT_DSN)%&queue_name=neuromancer'

 routing:

 App\Message\NeuromancerMessage: neuromancer

Retry

- In a large system, things will go wrong
 - Software crashes, machines reboot
 - Network issues
 - Service or system not available
 - Actual logic bugs
- Message goes back to queue
- Exponential backoff
- Failure transport

Keep workers running

- Long-running PHP processes
- Careful about memory leaks.
- supervisord or similar to restart
- Rejuvenation: Restart after number of messages, time or when memory consumption exceeds threshold
- Restart workers on deployment

Autoscaler

- Own service to start and stop workers
- Application registers its workers with autoscaler
- Autoscaler monitors queue sizes
- Spawns workers up to the limit when queue grows

What is going on?

- Copy all messages to an archive queue with a TTL
- Import state log table, track all steps

Message Transport

- Without a configured transport, messages are processed immediately
- Doctrine (using a database table)
- AMQP with ext-amqp (rabbitmq, kafka)
- Redis with ext-redis
- Amazon SQS
- Plug your own, e.g. to call a custom API

Did i hear <<microservices>>?



MICROSERVICES
FIX EVERYTHING

MICROSERVICES
ARE THE ROOT
OF ALL EVIL

Une Bataille, ~1750, François-Joseph Casanova

Messaging != microservices

- Our application has ~30 types of message and their processors
- Messages and processors all in the same application
- I would not want to maintain 30 tiny Symfony applications just for one single worker each
- Some separate applications that have nothing directly to do with the main data api



Thank you!

github.com/dbu/messenger-demo/

@dbu@phpc.social

David Buchmann, Liip SA