

A Tale of Two Cities: Blocking Code VS Non-Blocking Code

Bazlur Rahman
Staff Software Developer
DNAStack



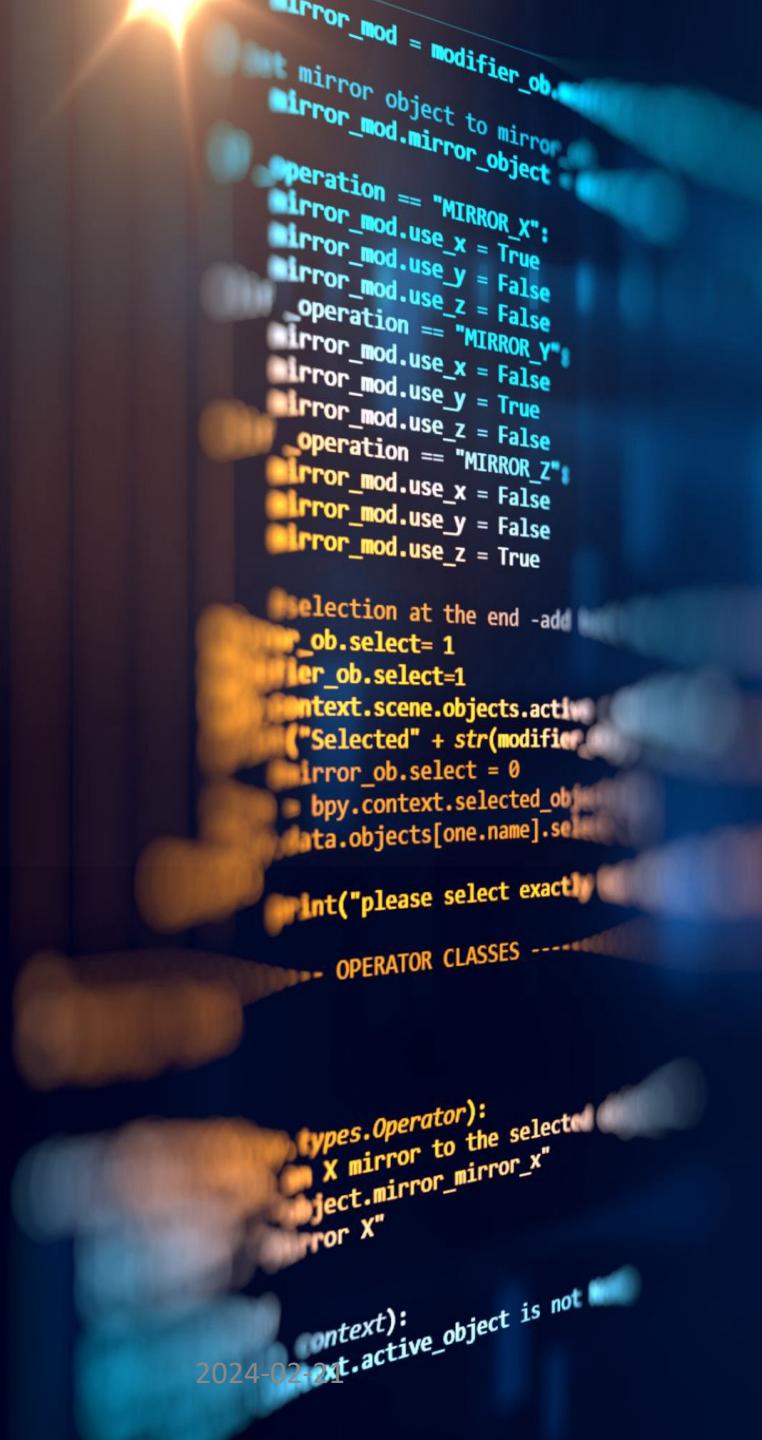
Agenda

2024-02-21

@bazlur_rahman

2





Threads—The Foundation of Java Concurrency

- Java is born with threads! When a Java program starts, the "main" thread is the birthplace of execution.
 - Threads give birth to threads - Method calls spawn execution within the caller's thread
 - Benefits we take for granted:
 - Structured control flow (think about how 'easy' writing sequential if/then/else logic is...)
 - Local variables for methods
 - Stack traces to solve problems
 - Ability to schedule work across CPUs

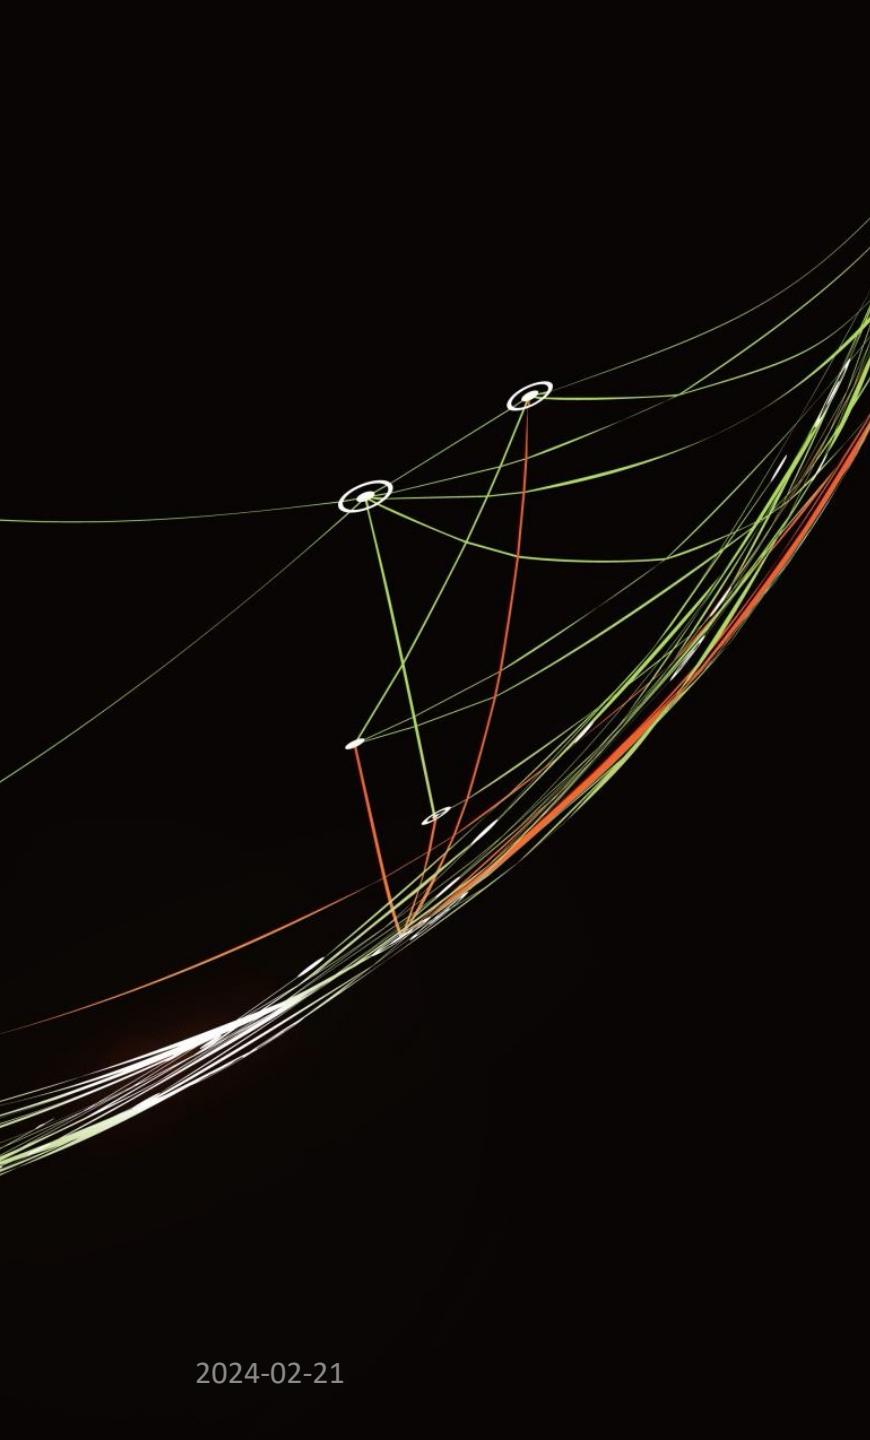
The screenshot shows the IntelliJ IDEA interface with a Java project named "tale-of-two-cities". The project structure on the left includes ".idea", "out", "src", and several source files like "CountThreads.java", "Main.java", and "Playground.java". The "Main.java" file is open in the editor, displaying a simple "Hello world!" application.

In the bottom panel, the "Debug" tool window is active, showing a thread dump titled "ca.bazlur.jcon2023.Main". The dump lists various threads and their states:

- "main"@... RUNNING
- "main"@1 in group "main": RUNNING
- "Common-Cleaner"@770 in group "InnocuousThreadGroup": WAIT
- "Finalizer"@768: WAIT
- "Notification Thread"@732: RUNNING
- "Reference Handler"@767: RUNNING
- "Signal Dispatcher"@769: RUNNING

The status bar at the bottom indicates the build configuration is "ca.bazlur.jcon2023.Main", and the file encoding is "UTF-8".

```
void main() throws InterruptedException {  
    var thread = Thread.ofPlatform().start(() ->  
        System.out.println(""  
            + "Hello from the brand new thread!"  
            + "I was just born, but I could already go for a cup of coffee!"  
            + """));  
    System.out.println("Main thread: Time to fetch some coffee!");  
    thread.join();  
}
```



Java Threads: A Wrapper Around OS Threads

- Java threads rely on operating system threads.
- OS-managed threads are often called "**Platform threads.**"
- Pre-JDK 21 had a one-to-one correspondence between Java threads and OS threads.
- This approach could limit performance.



The Cost of Platform Threads

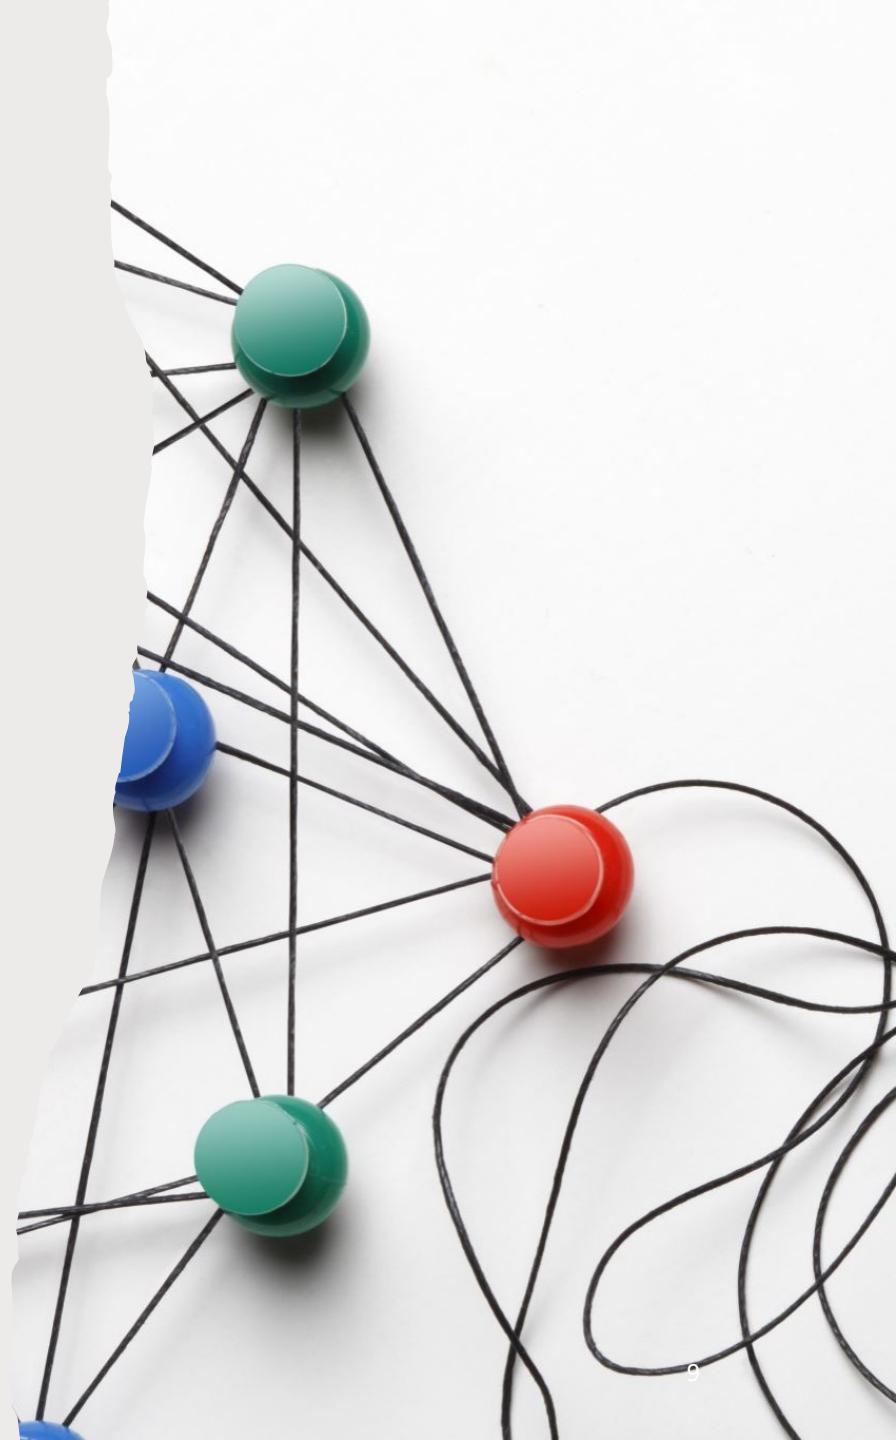
- OS thread creation takes time and significant resources.
- Large "pre-commitment" needed for stacks
 - often megabytes in size (can sometimes be tuned, but with risks).
- This introduces a practical cap on the number of concurrent threads in an application.

How Many Platform Threads?

```
void main() {  
    var counter = new AtomicInteger();  
    System.out.println("Welcome to the Thread Race Marathon!");  
  
    while (true) {  
        Thread.ofPlatform().start(() -> {  
            int count = counter.incrementAndGet();  
            if (count % 1000 == 0) {  
                System.out.println(STR."We have created \{count} threads so far. Keep  
going!");  
            }  
            LockSupport.park(); // Suspend the thread, simulating the racer taking a break  
        });  
    }  
}
```

The Challenges of Shared-State Concurrency

- While invaluable, threads become challenging when we manipulate shared data across threads.
- Issues to keep in mind:
 - Memory visibility (ensuring one thread sees another's changes)
 - Safeguarding against race conditions and data corruption
 - Mastering locks and synchronization for protection.

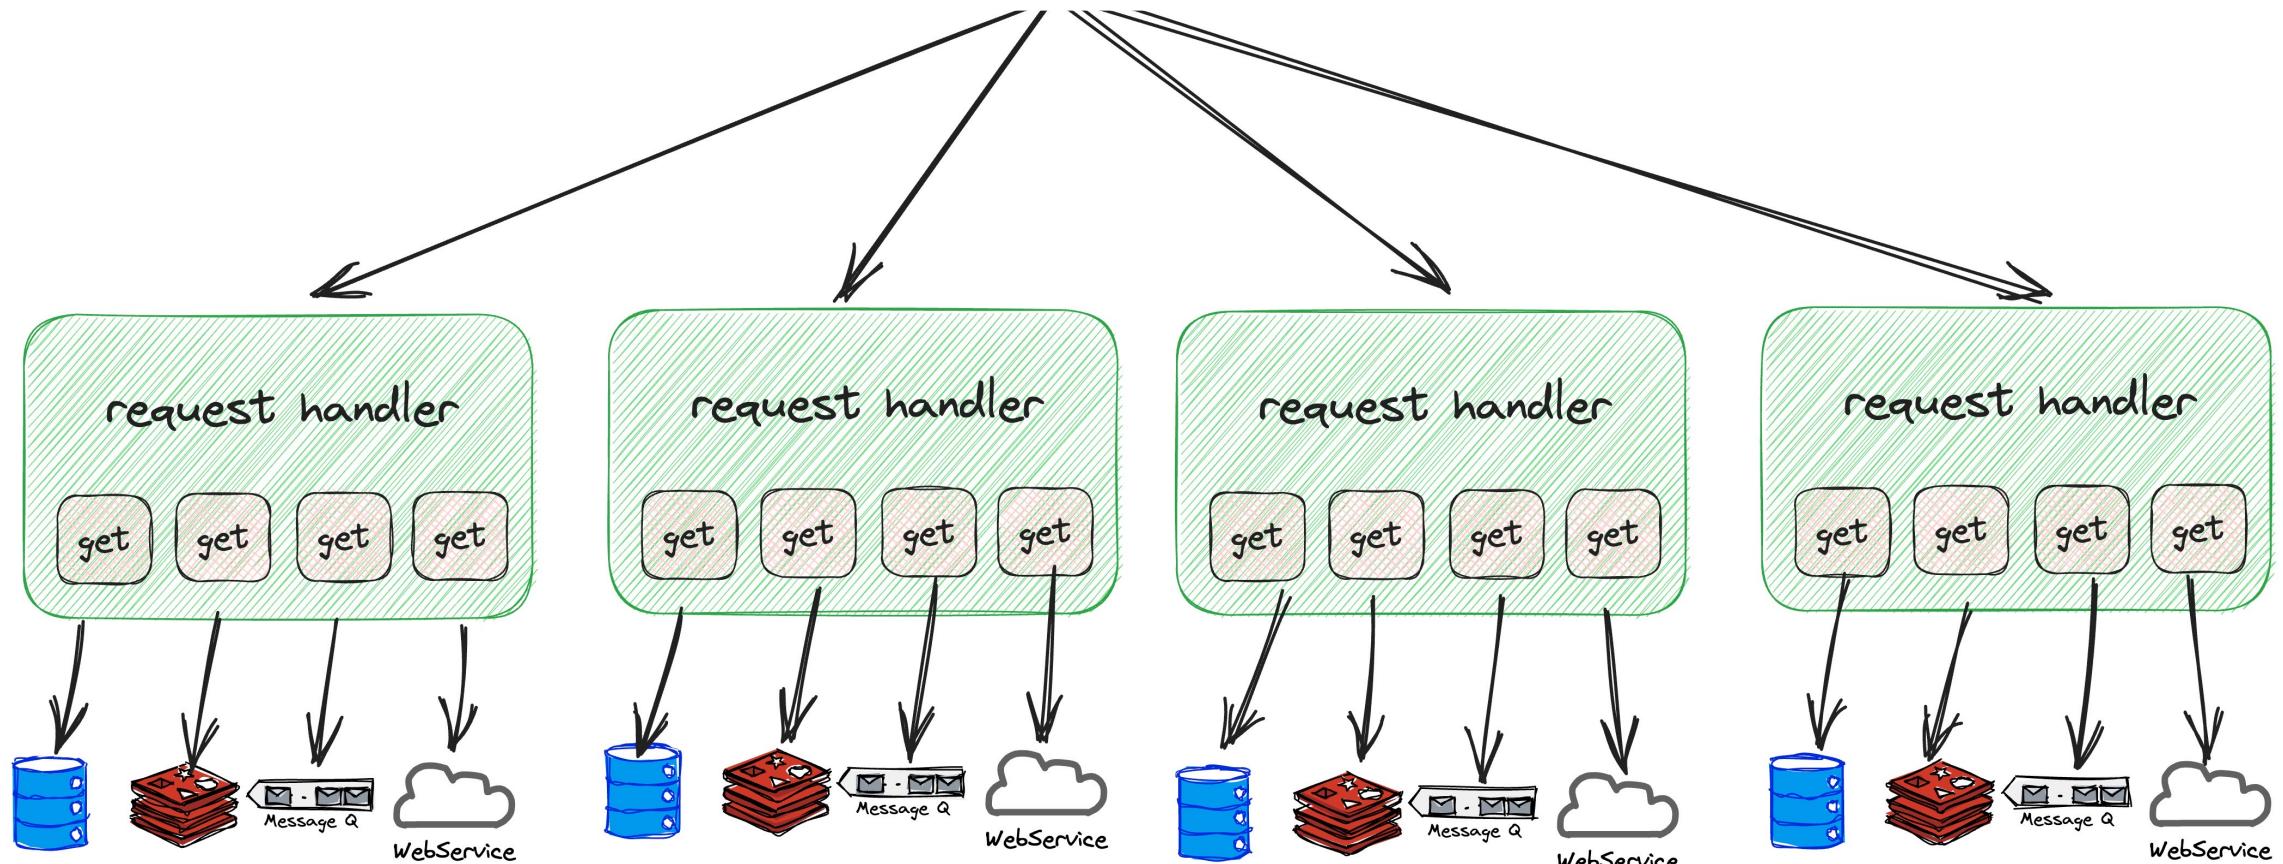


Beyond the Basics

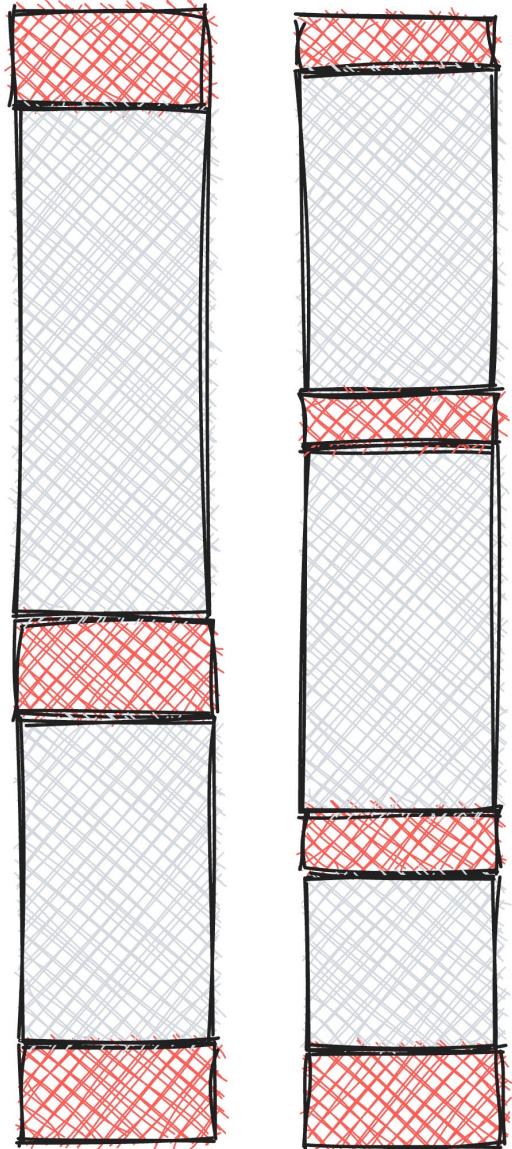
- Despite thread complexities, a huge amount of what we accomplish day-to-day relies on them working silently...
- Java concurrency is rich - we only highlight a part of the landscape today.
- Our focus today... how do we manage tasks that might take a while without tying up our precious threads?

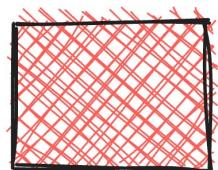


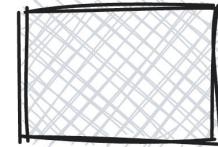
Common Server Fanout

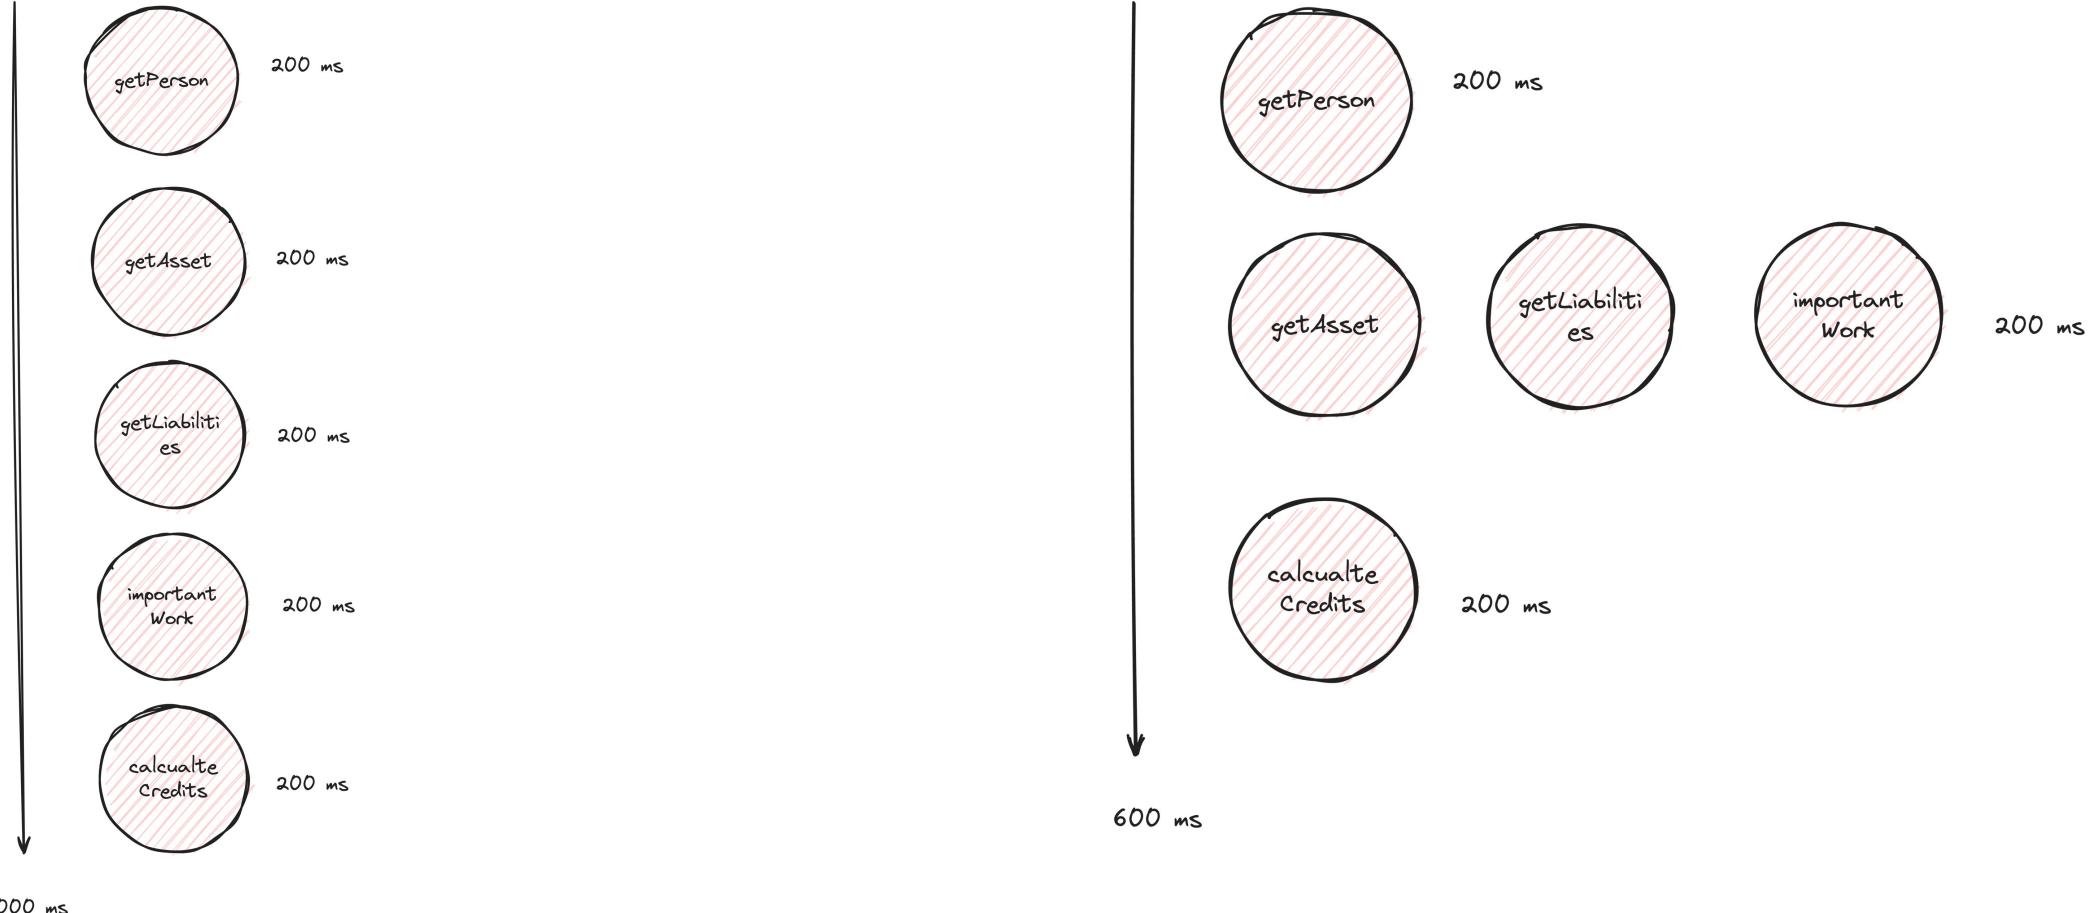


Outgoing requests



 - processing

 - I/O



Classical Implementation

```
var atomicReference = new AtomicReference<Person>();  
  
Thread.ofPlatform().start(() -> {  
    var person = getPerson(1L);  
    atomicReference.set(person);  
});
```

```
Credit calculateCreditClassicalThreading(Long personId) throws InterruptedException {
    var person :Person = getPerson(personId);

    var assetsWrapper = new AtomicReference<List<Asset>>();
    var assetsThread :Thread = Thread.ofPlatform().unstarted(() -> {
        List<Asset> assets = getAssets(person);
        assetsWrapper.set(assets);
    });

    var liabilitiesWrapper = new AtomicReference<List<Liability>>();
    var liabilitiesThread :Thread = Thread.ofPlatform().unstarted(() -> {
        List<Liability> liabilities = getLiabilities(person);
        liabilitiesWrapper.set(liabilities);
    });

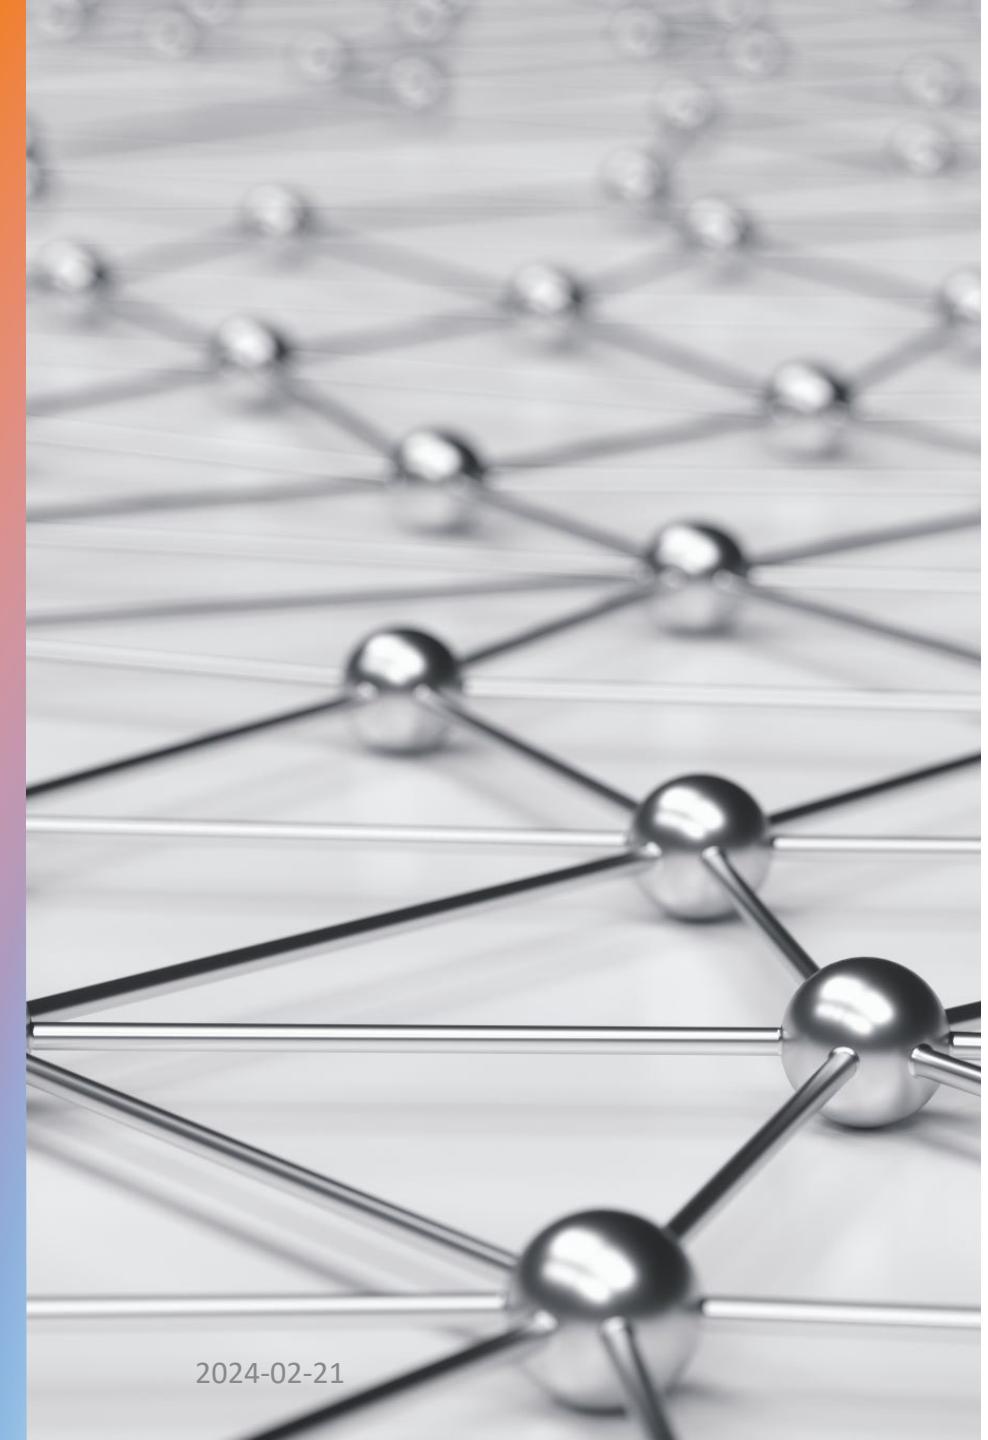
    var importantWorkThread :Thread = Thread.ofPlatform().unstarted(this::importantWork);

    assetsThread.start();
    liabilitiesThread.start();
    importantWorkThread.start();

    assetsThread.join();
    liabilitiesThread.join();

    var credit :Credit = calculateCredits(assetsWrapper.get(), liabilitiesWrapper.get());

    importantWorkThread.join();
    return credit;
}
```



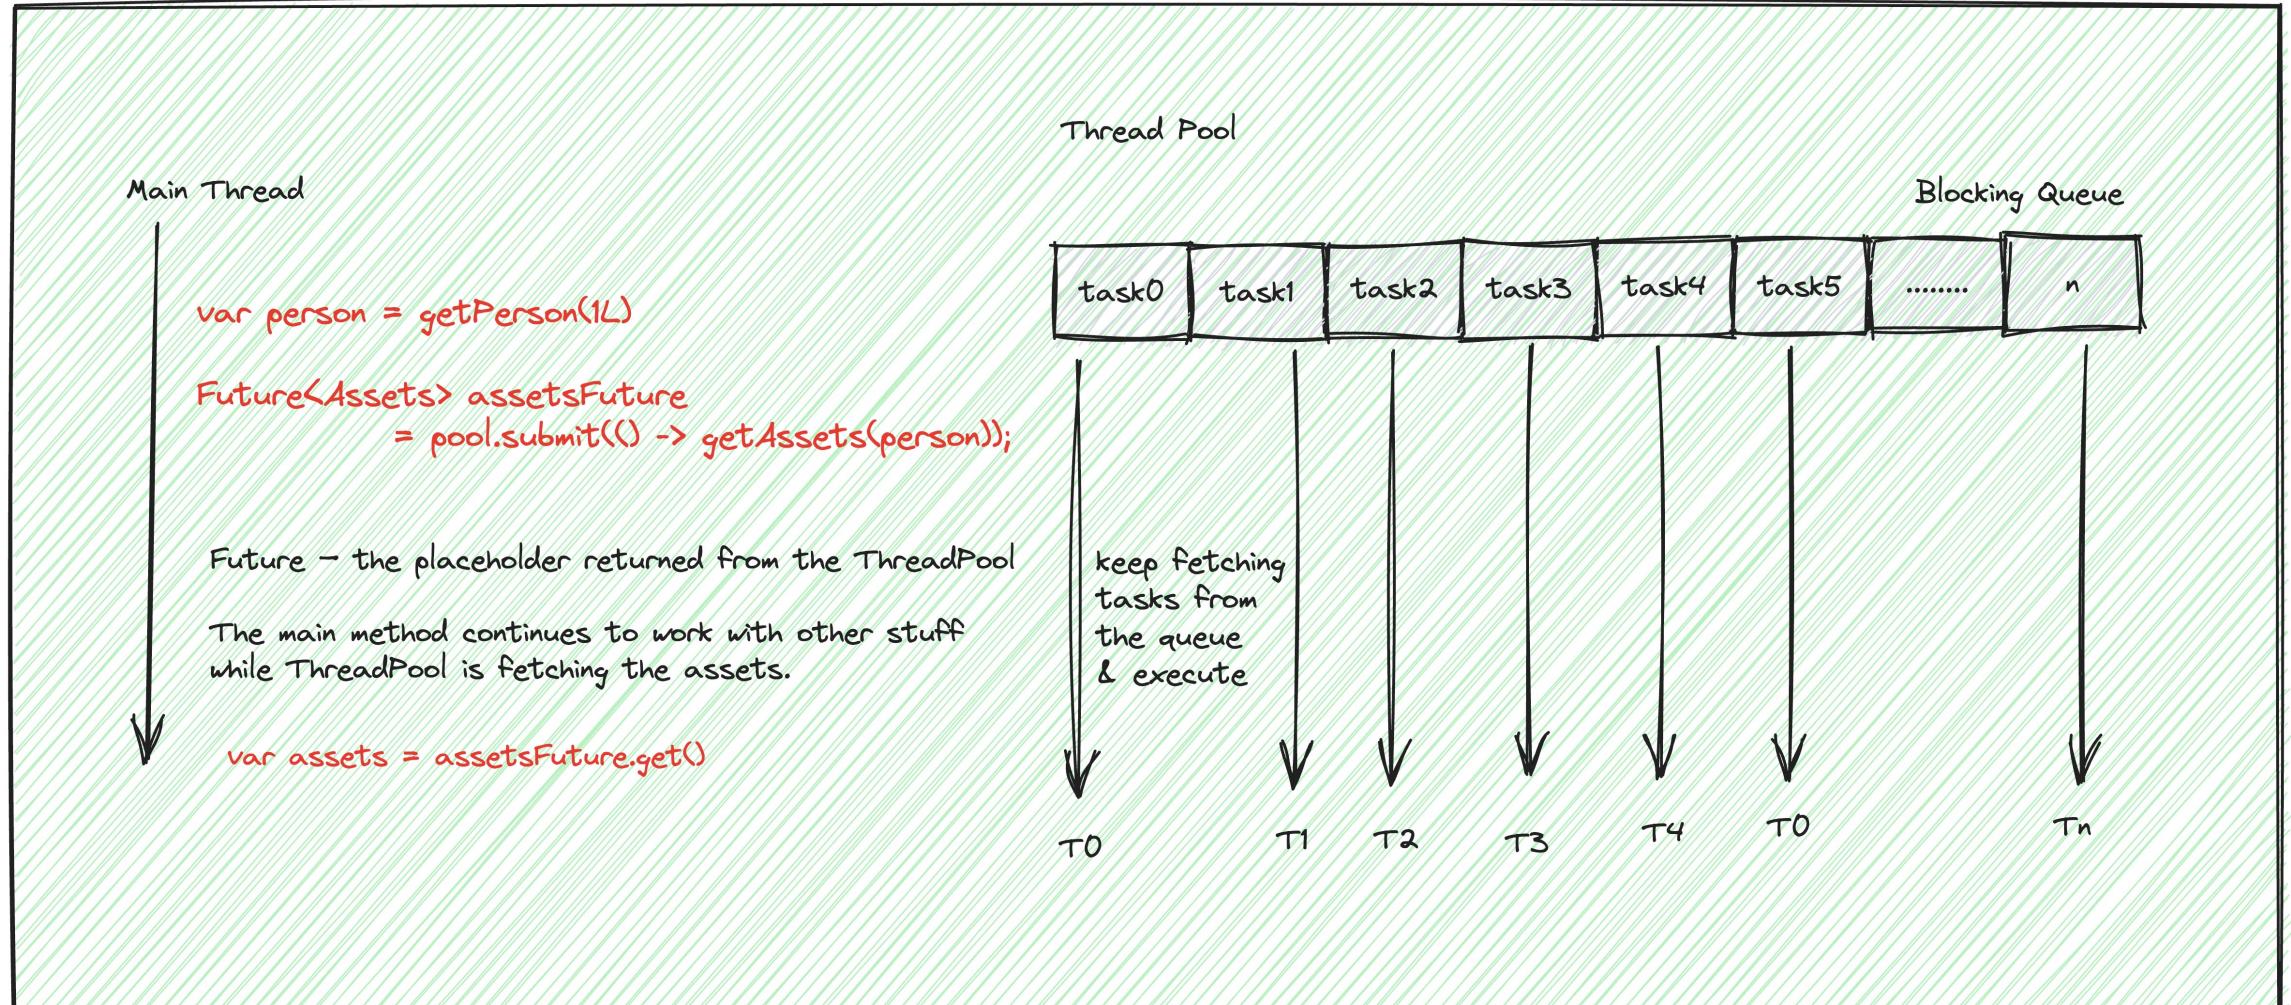
The Inception – Java's Early Concurrency Tools

- Future/Callable
 - This means to execute tasks asynchronously
 - The Future object represents the eventual result of the computation
 - The Callable interface provides a way to define the work to be done

```
Credit calculateCreditWithExecutor(Long personId) throws ExecutionException, InterruptedException {
    var person :Person = getPerson(personId);
    var assetsFuture :Future<List<...>> = executor.submit(() -> getAssets(person));
    var liabilitiesFuture :Future<List<...>> = executor.submit(() -> getLiabilities(person));
    executor.submit(this::importantWork);

    return calculateCredits(assetsFuture.get(), liabilitiesFuture.get());
}
```

How Traditional Thread Pool Works



Challenges with Executors and Future/Callable

- **Awkward Composability:** Chaining dependent asynchronous operations leads to cumbersome nesting.
- **Blocking Future.get():** Retrieving results can potentially block threads, negating some non-blocking gains.
- **Limited Error Handling:** Propagating exceptions across asynchronous stages can be tricky to manage.
- **Executor Management:** Tuning thread pools for optimal performance requires careful consideration.

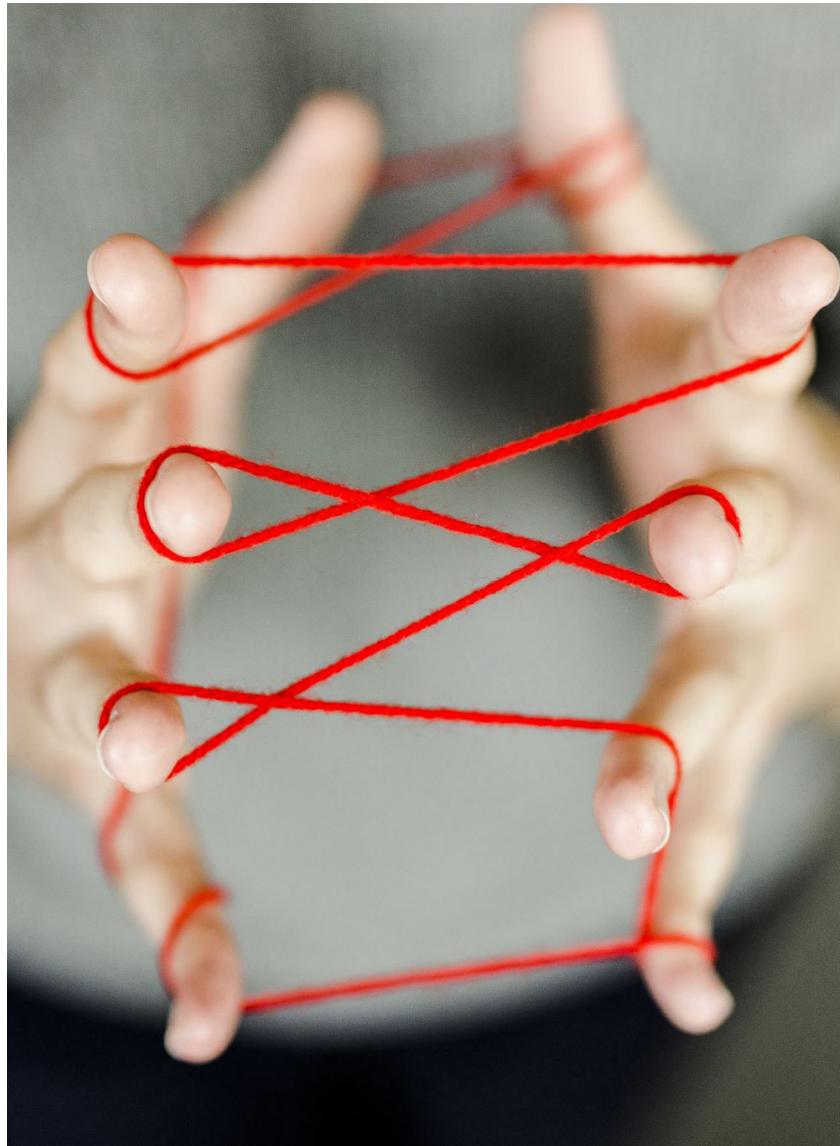


CompletableFuture to the Rescue

- Methods like `thenApply`, `thenCombine`, and `handle` enable fluent chaining of asynchronous operations without awkward nesting.
- The chaining pattern remains non-blocking; subsequent dependent tasks are asynchronously queued.
- Fine-grained control with `exceptionally` and similar methods provides more robust exception management across asynchronous tasks.
- While granular control is possible, many common operations have defaults, making it less of a required upfront burden.
- The chaining style starts to bridge non-blocking execution with syntax that feels closer to sequential code for many developers.



```
Void calculateCreditWithCompletableFuture(Long personId) {
    return CompletableFuture.supplyAsync(() -> getPerson(personId)).thenComposeAsync(person ->
        CompletableFuture.allOf(
            CompletableFuture.runAsync(this::importantWork),
            CompletableFuture.supplyAsync(() -> getAssets(person))
                .thenAcceptBothAsync(CompletableFuture.supplyAsync(() -> getLiabilities(person)),
                    this::calculateCredits)))
    .join();
}
```

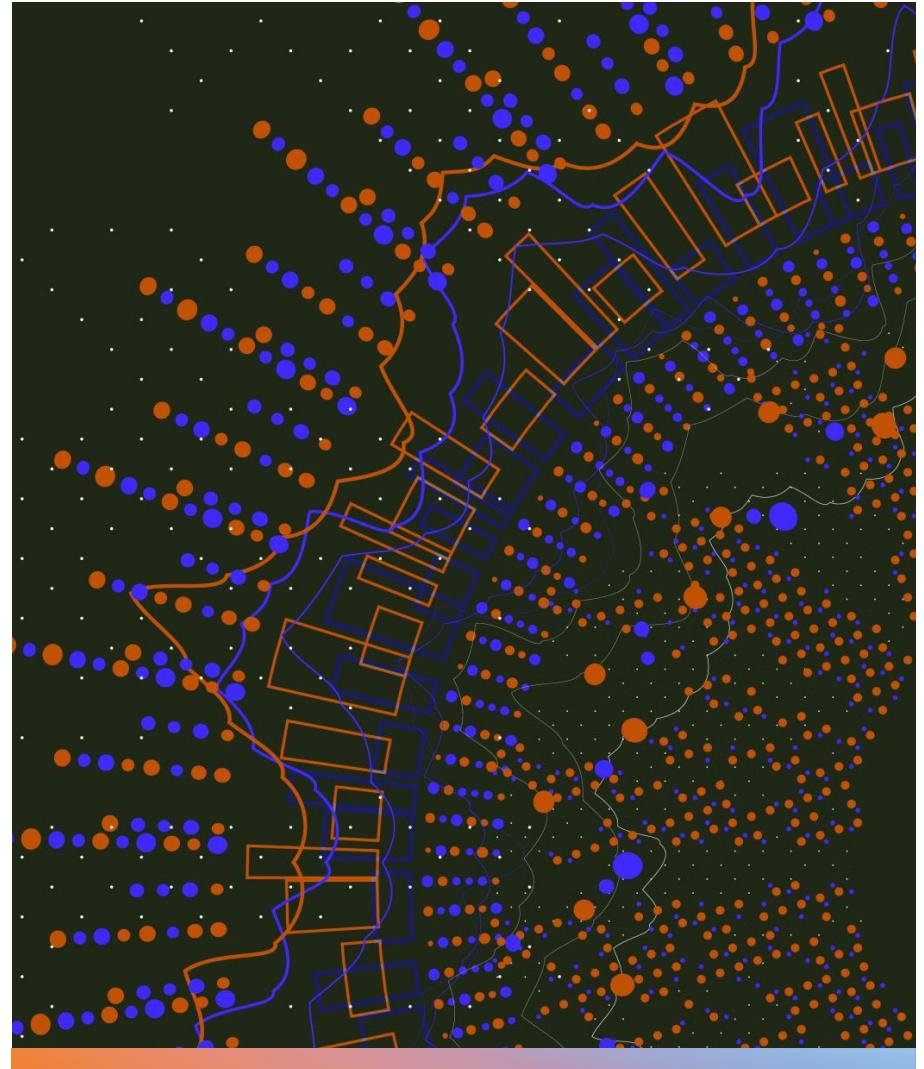


The Limitations of CompletableFuture

- **Complexity Can Creep Back In:** Intricate logic leads to dense chains, readability declines.
- **Control-Flow Limits:** Non-linear logic or complex task coordination becomes awkward.
- **Imperative Style Persists:** We still 'dictate' how tasks connect, not their data dependencies.
- **Debugging Challenges:** Async stacks remain tricky; traditional tools become less effective.

The Reactive Stack – Benefits

- **Flow Control (Backpressure):** Publishers and subscribers negotiate data rates, preventing components from being overwhelmed.
- **Declarative Composability:** Build complex stream transformations through chaining expressive operators (e.g. filter, map, combine).
- **Asynchronous by Design:** Optimized for concurrency and handling events that come at unpredictable intervals.
- **Resilience:** Gracefully handles errors and inconsistencies within data streams.
- **Popular Framework:** RxJava, Akka, Eclipse Vert.x , Spring WebFlux, Slick



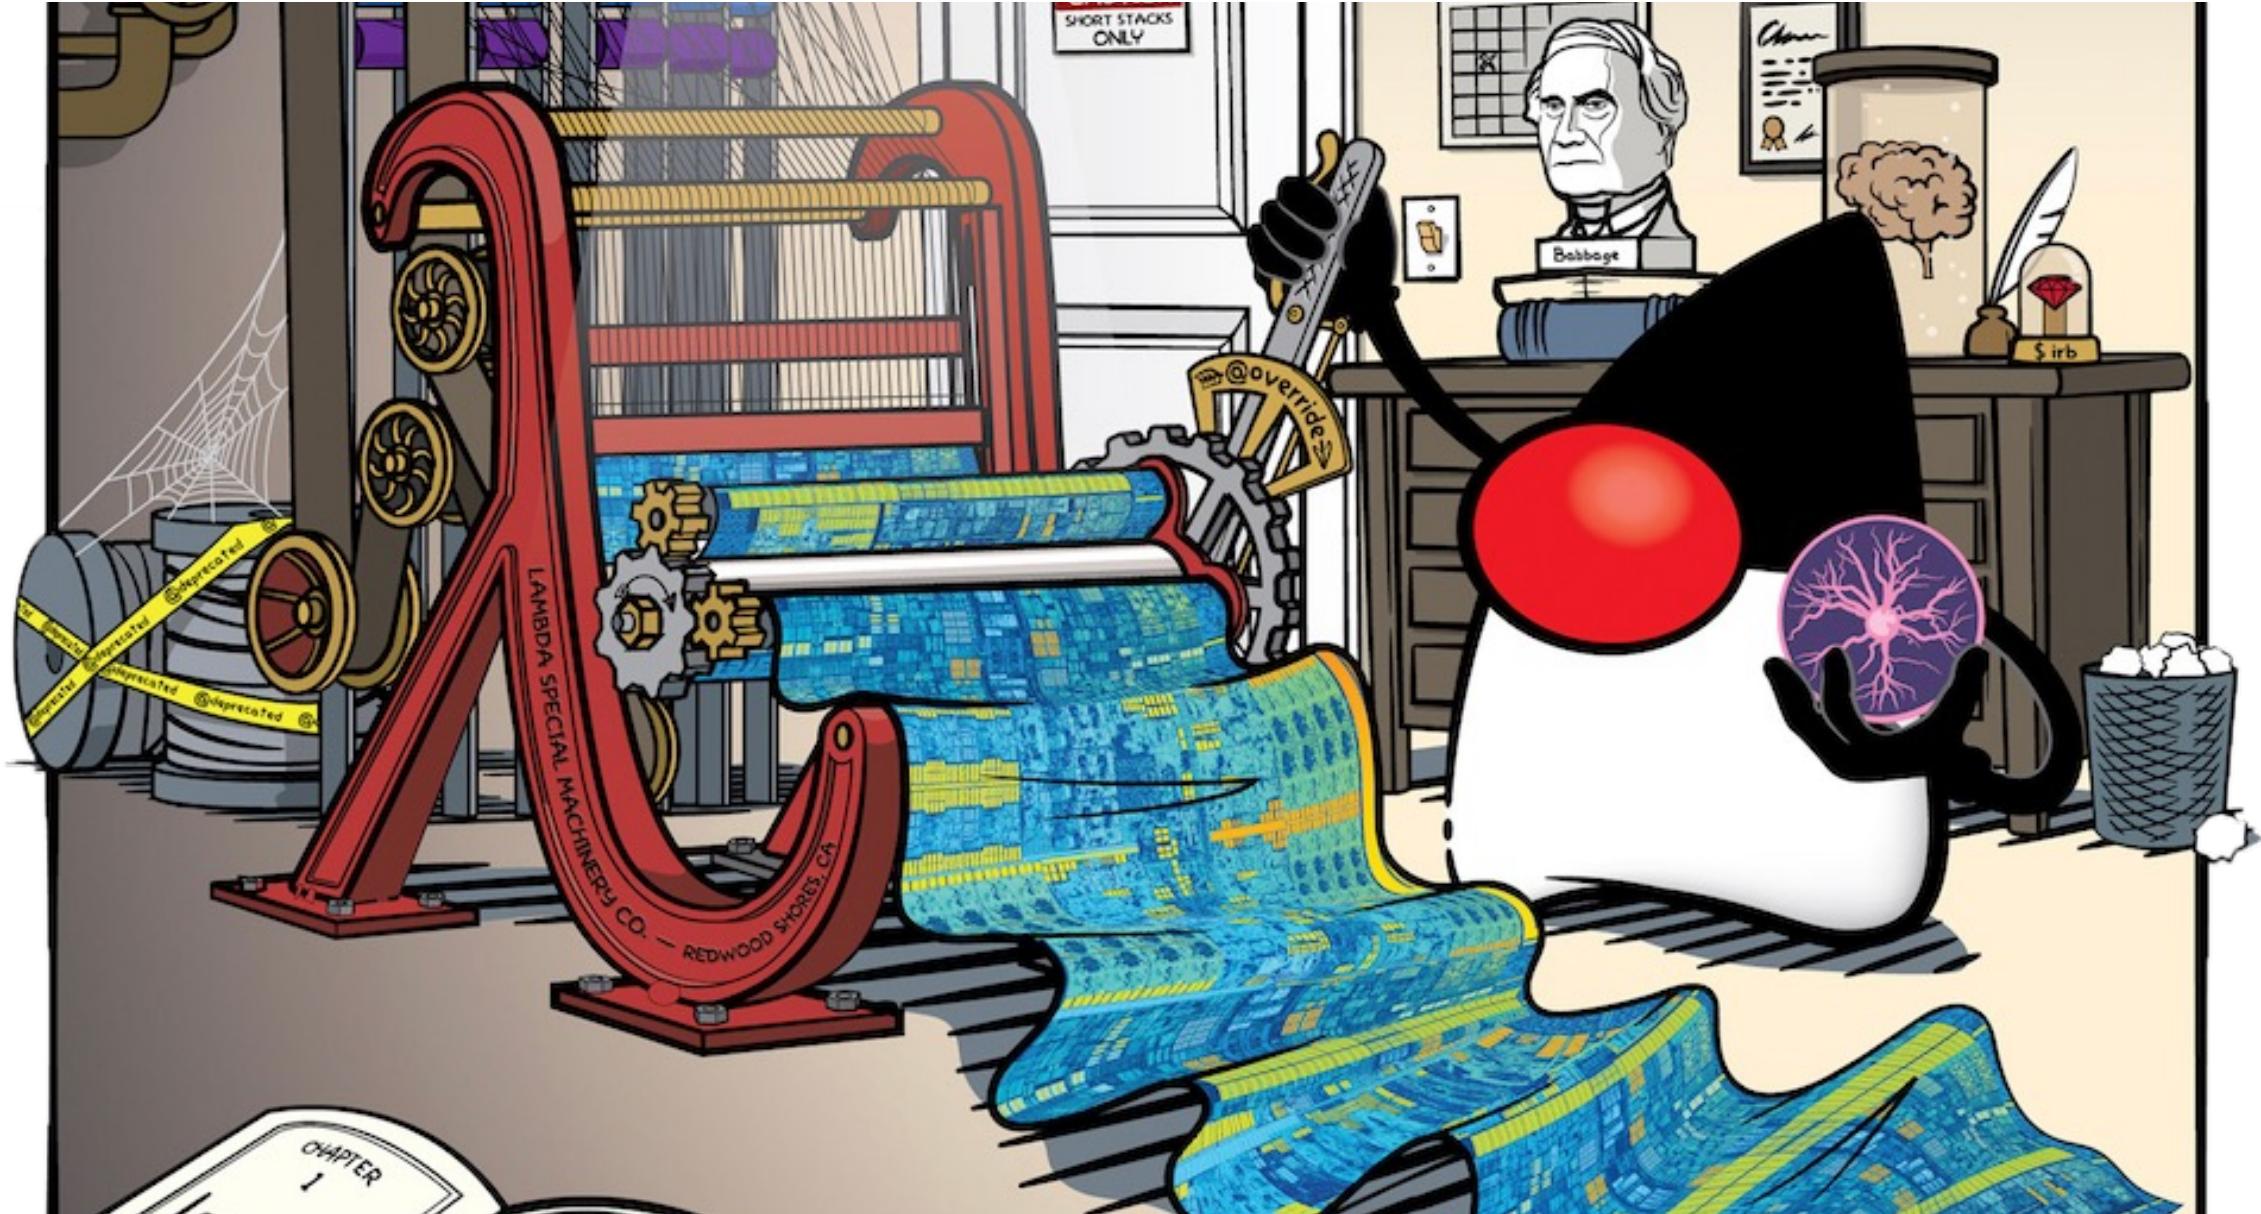
The Reactive Stack – Considerations

- **Learning Curve:** The shift to data-flow thinking involves mastering new concepts and patterns.
- **Debugging:** Asynchronous data-flows can pose unique debugging challenges.
- **Code Readability:** Complex transformations may benefit from careful documentation for maintainability.
- **When to Choose:** Best suited for highly asynchronous, event-driven scenarios or problems that map naturally to a data-stream processing model.



```
import io.reactivex.rxjava3.core.Single;

Single<Credit> calculateCredit (Long personId) { return
Single.fromCallable(this::importantWork)
    .flatMap(aVoid -> Single.fromCallable(() -> getPerson(personId))
        .flatMap(person -> Single.zip(
            Single.fromCallable(() -> getAssets(person)), Single.fromCallable(() ->
getLiabilities(person)),
this::calculateCredits
        )
    )
);
}
```

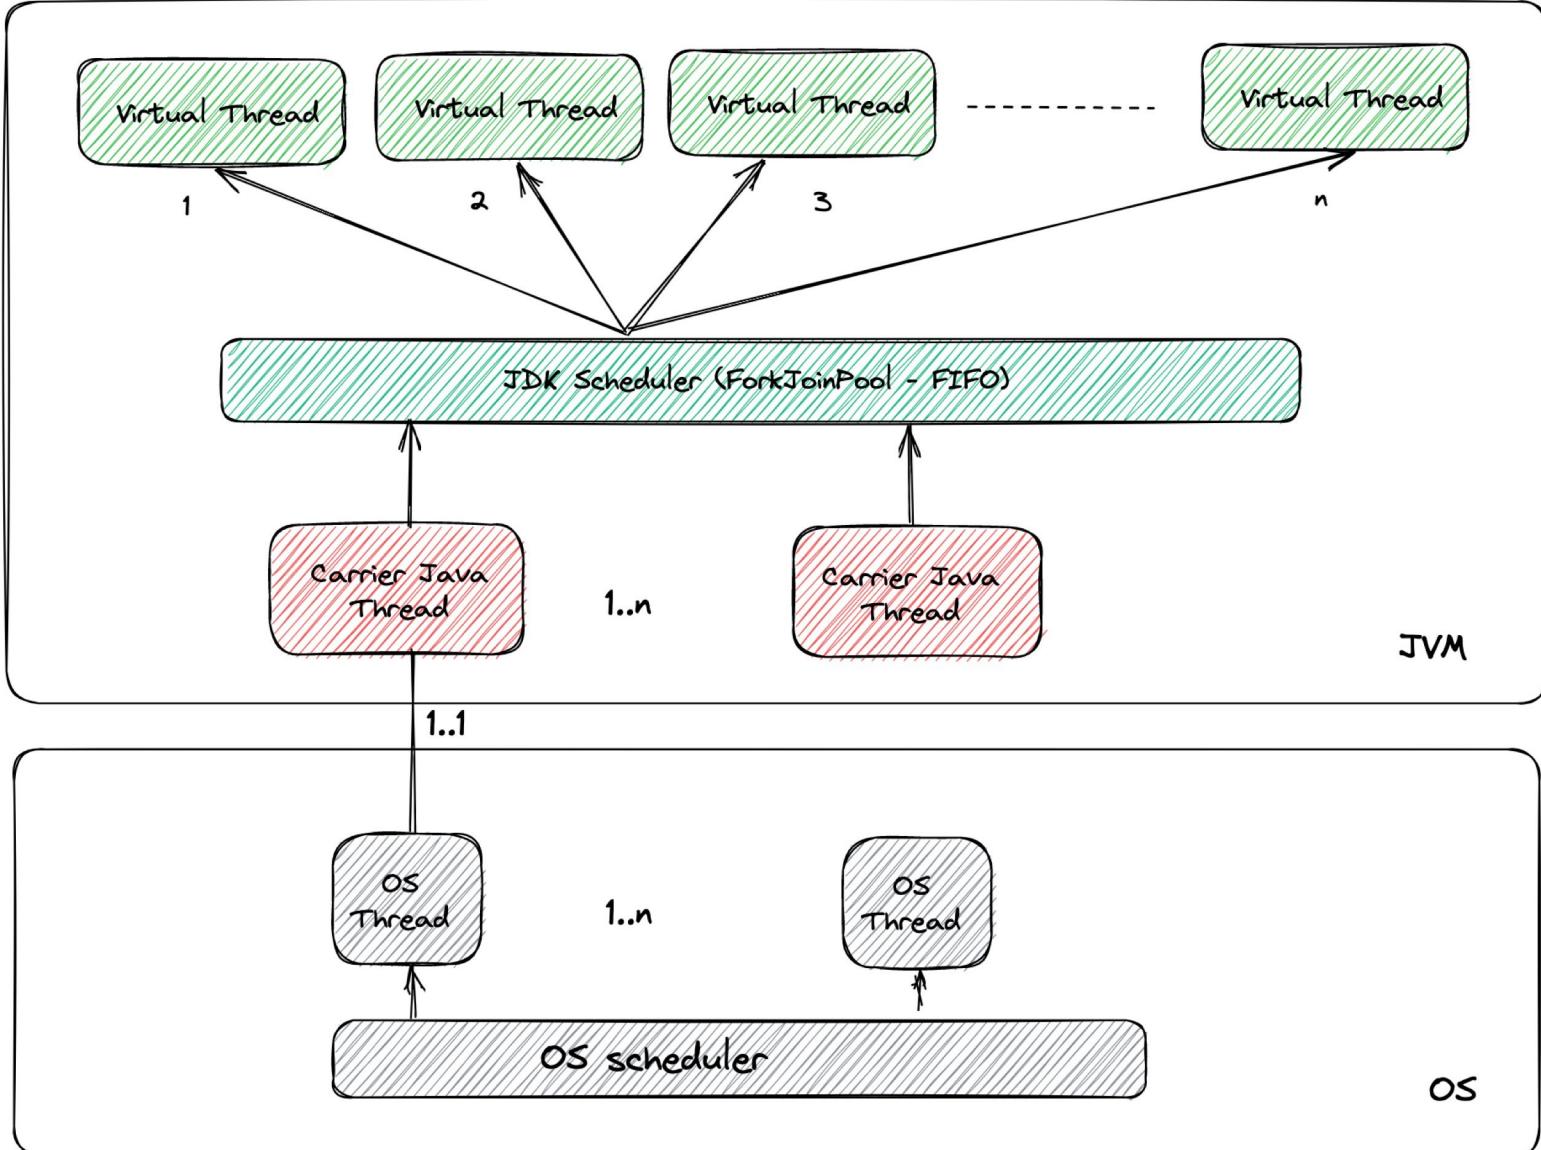




Demo

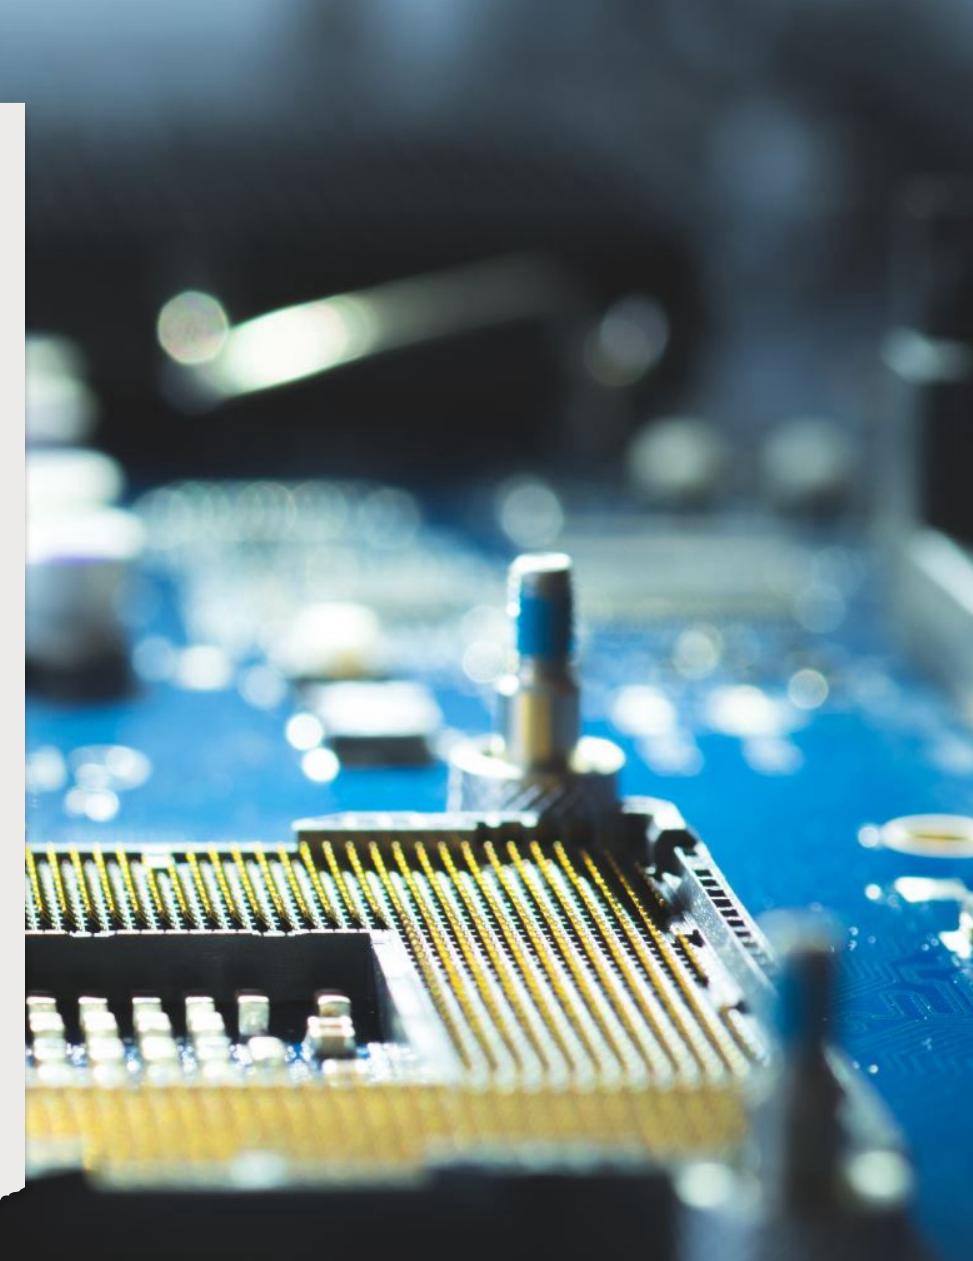
Thread Evolution - Virtual Threads vs. Platform Threads

Feature	OS-Managed Threads	Virtual Threads
Management	OS-managed, heavier resource footprint	JVM-managed, lightweight with dynamic memory
Quantity	Limited in number by memory constraints	Potential for millions of virtual threads
Context Switching Overhead	Context switching can be more overhead	Seamless mounting/unmounting reduces overhead



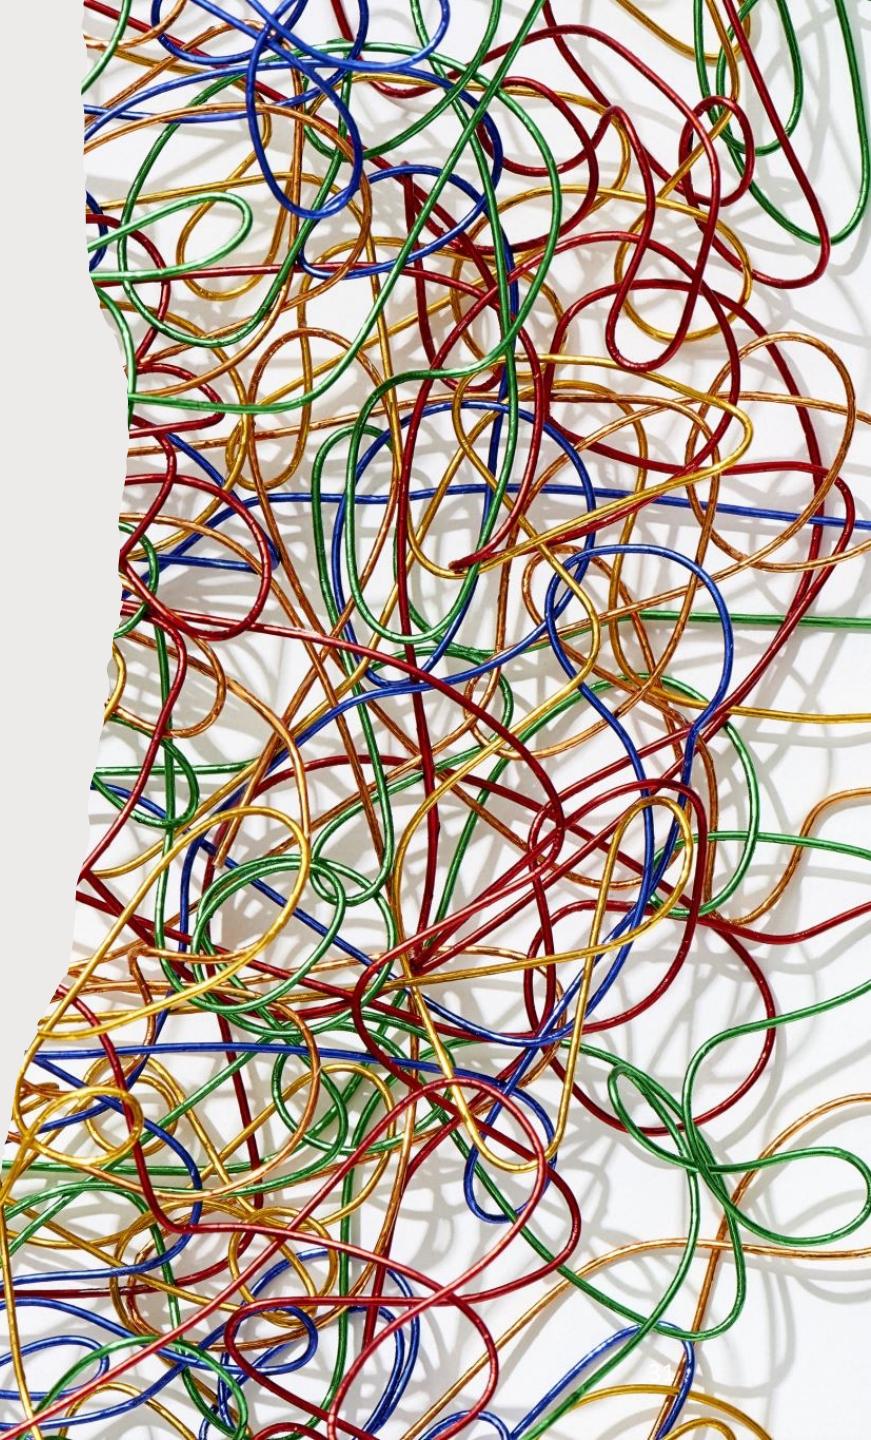
Benefits

- Virtual threads don't consume CPU while they are waiting/sleeping
- Technically, we can have millions of virtual threads
- Threads are cheap
- The request-per-thread style of server-side programming becomes as manageable again as before.
- Improved throughput
- On Blocking I/O, virtual threads get automatically suspended
- No longer need to maintain ThreadPool



It's About Scalability!

- The Bottleneck: Traditional Threads and I/O
 - Many server tasks involve waiting (network responses, database queries).
 - Platform threads become idle during waits, wasting a precious resource.
 - More platform threads = bigger memory overhead, even if they just sit inactive!
- Virtual Threads: Breaking the Constraint
 - Minimal footprint allows MANY virtual threads, most not consuming a platform thread at any given moment.
 - During I/O waits, virtual threads 'unmount' from their carriers, staying efficient with heap footprint instead.
 - Carriers dynamically pick up other ready-to-run virtual threads – maximizing work done.
- Hardware's True Power Realized
 - Before, we couldn't fully use modern CPUs due to the artificial limit on threads.
 - Virtual threads let applications squeeze the maximum work out of available cores.
 - This doesn't make INDIVIDUAL tasks finish faster, it keeps everything flowing at top potential

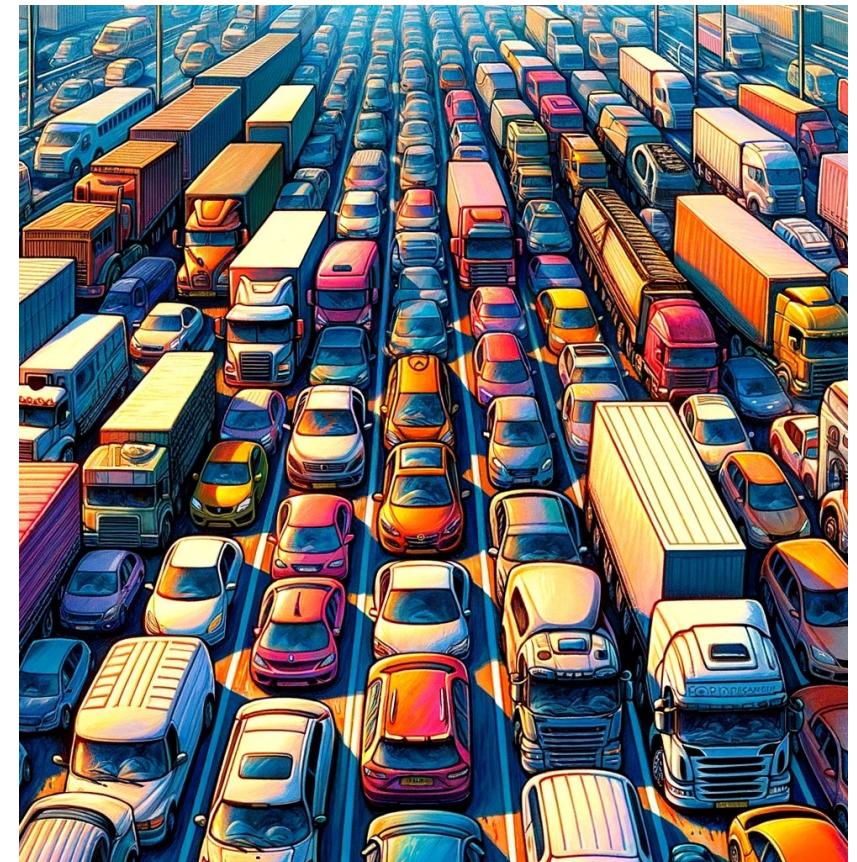


Use Semaphores to Limit Concurrency

```
Semaphore sem = new Semaphore(10);  
...  
Result foo() {  
    sem.acquire();  
    try {  
        return callLimitedService();  
    } finally {  
        sem.release();  
    }  
}
```

The Challenge of Pinning

- Virtual threads map to carrier threads (platform threads) for real work to happen.
- Blocking operations (old-school sync, some native code) "pin" a virtual thread to its carrier.
- A pinned thread stalls the carrier, holding up other virtual threads waiting in line.
- Excessive pinning in high-concurrency systems negates some scalability gains.



Solution: Catch 'em in the Act.

-Djdk.tracePinnedThreads=full

Stacktrace

```
...
...
...
java.base/sun.nio.cs.StreamDecoder.close(StreamDecoder.java:249)
    java.base/java.io.InputStreamReader.close(InputStreamReader.java:203)
    okhttp3.ResponseBody$BomAwareReader.close(ResponseBody.kt:217)
    java.base/java.io.BufferedReader.implClose(BufferedReader.java:636)
java.base/java.io.BufferedReader.close(BufferedReader.java:627) <== monitors:1
    com.fasterxml.jackson.core.json.ReaderBasedJsonParser.closeInput(ReaderBasedJsonParser.java:235)
    com.fasterxml.jackson.core.base.ParserBase.close(ParserBase.java:392)
    com.fasterxml.jackson.databind.ObjectMapper._readMapAndClose(ObjectMapper.java:4833)
```

How Long Was It Stuck?

- Find your application's PID with the **jps** command.
- Record some data:

```
jcmd <PID> JFR.start duration=200s filename=myrecording.jfr
```

```
jfr print --events jdk.VirtualThreadPinned myrecording.jfr
```

```
jdk.VirtualThreadPinned {
    startTime = 13:21:28.226 (2024-02-12)
duration = 108 ms
    eventThread = "tomcat-handler-8" (javaThreadId = 76, virtual)
    stackTrace = [
        java.lang.VirtualThread.parkOnCarrierThread(boolean, long) line: 677
        java.lang.VirtualThread.parkNanos(long) line: 636
        java.lang.System$2.parkVirtualThread(long) line: 2648
        jdk.internal.misc.VirtualThreads.park(long) line: 67
        java.util.concurrent.locks.LockSupport.parkNanos(long) line: 408
        ...
    ]
}

jdk.VirtualThreadPinned {
    startTime = 13:21:40.117 (2024-02-12)
duration = 93.2 ms
    eventThread = "tomcat-handler-26" (javaThreadId = 94, virtual)
    stackTrace = [
        java.lang.VirtualThread.parkOnCarrierThread(boolean, long) line: 677
        java.lang.VirtualThread.parkNanos(long) line: 636
        java.lang.System$2.parkVirtualThread(long) line: 2648
        jdk.internal.misc.VirtualThreads.park(long) line: 67
        java.util.concurrent.locks.LockSupport.parkNanos(long) line: 408
        ...
    ]
}
```

About Me

- Staff Software Developer
- Java Champion
- Jakarta EE Ambassador
- JUG Leader
- Published Author
- InfoQ Editor of Java Queue
- Editor of Foojay.io



DNA STACK

A screenshot of an email newsletter titled "The Coding Café". It features a thumbnail of a steaming coffee cup, the author's photo (A N M Bazlur Rahman), and a brief description: "Introducing 'The Coding Café' – your go-to source for the perfect blend of Java programming and software development". It is described as "Published monthly".

The Coding Café
Introducing "The Coding Café" – your go-to source for the perfect blend of Java programming and software development
By A N M Bazlur Rahman
Java Champion | Senior Software Engineer | JUG...
Published monthly

https://twitter.com/bazlur_rahman
<https://www.linkedin.com/in/bazlur/>
<https://foojay.io/today/author/bazlur-rahman/>
<https://bazlur.ca>

